

# Deep Learning

Transformer

Lionel Fillatre

2025-2026


# Outline

- Attention Mechanism
- Single-Head Attention
- Multi-Head Attention
- Decode Only Transformer
- Input Encoding
- How does training work?
- Encoder-Decoder Transformer
- Conclusion

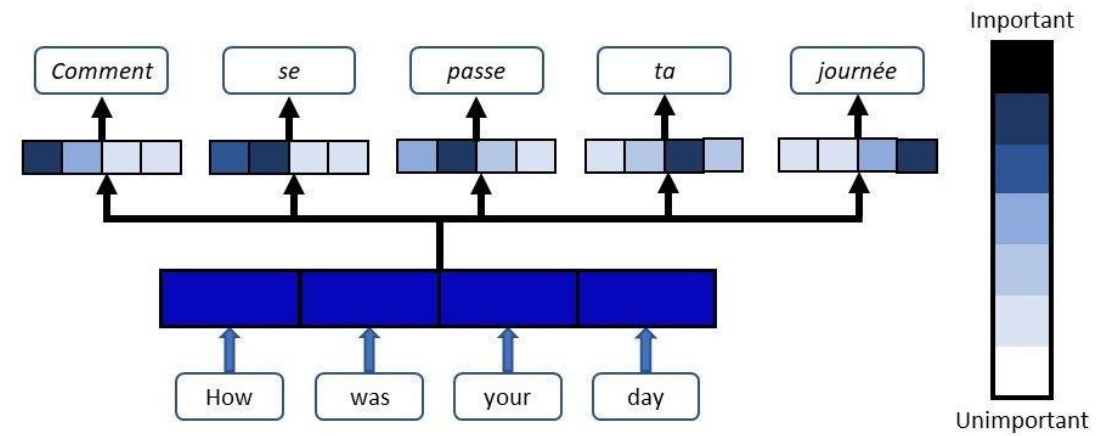
# Attention Mechanism

# Reminder: Encoder/Decoder

- The encoder-decoder model, translating the sentence “she is eating a green apple” to Chinese.
- The visualization of both encoder and decoder is unrolled in time.

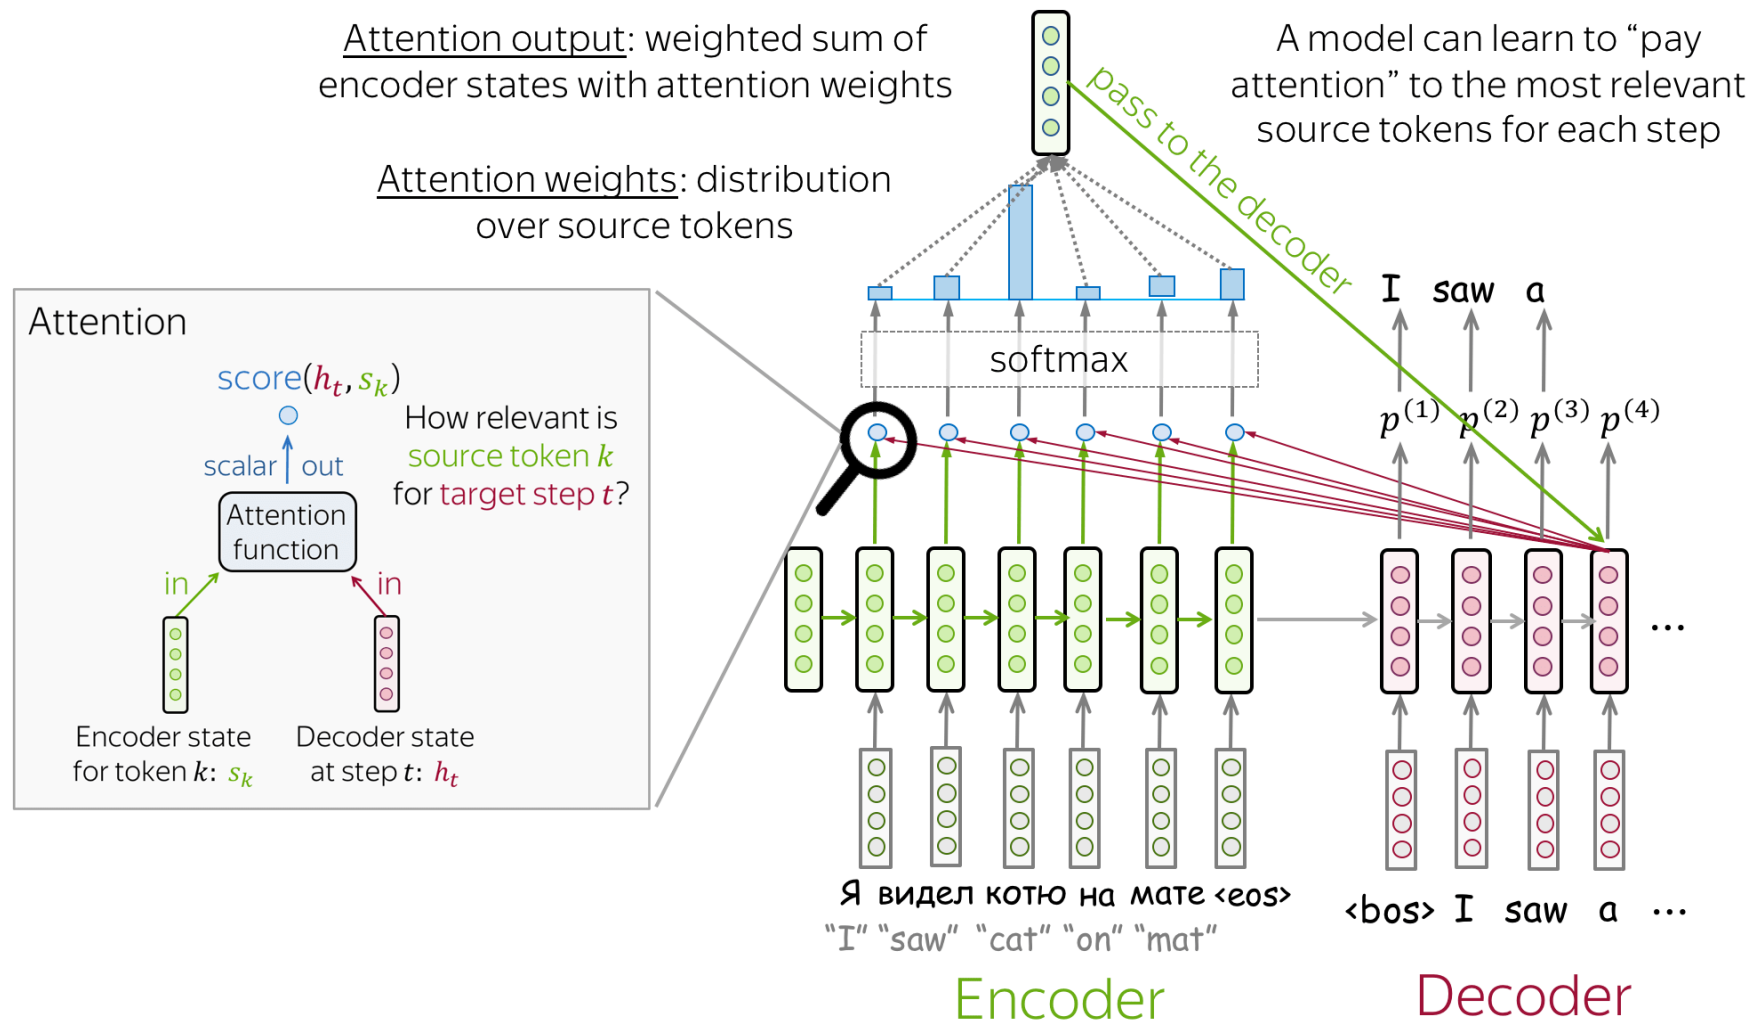


# Decoder with attention

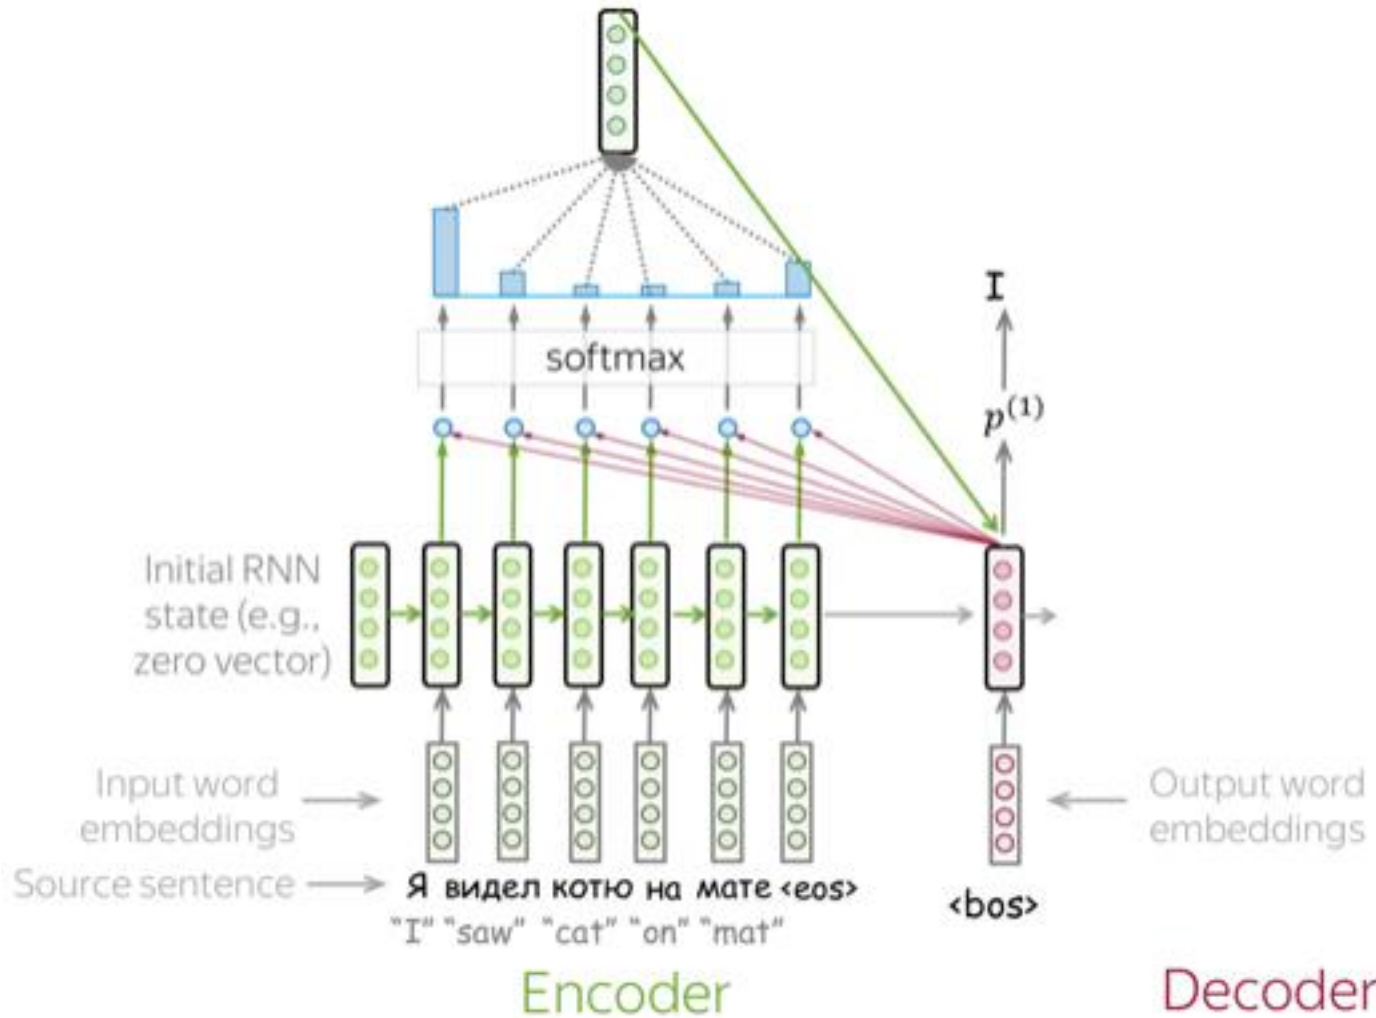


- The decoder is still a RNN
  - RNN decoder hidden state:  $s_{t'} = f(y_{t'-1}, s_{t'-1}, c_{t'})$
- With an RNN, each conditional probability is modeled as
$$p(y_{t'} | y_1, \dots, y_{t'-1}, x) = g(y_{t'-1}, s_{t'}, c_{t'})$$
  - $g$  is a nonlinear, potentially multi-layered, function that outputs the probability of  $y_{t'}$
- Principle: the probability is conditioned on a distinct context vector  $c_{t'}$  for each target word  $y_{t'}$


# An illustration



# An illustration




# An illustration







# An illustration




# An illustration




# An illustration



# An illustration




# An illustration





# An illustration



# Annotation

- The context vector  $c_i$  depends on a sequence of annotations  $h = (h_1, h_2, \dots, h_T)$  to which an encoder maps the input sentence.
- Each annotation  $h_i$  contains information about the whole input sequence (especially for a bidirectional RNN) with a strong focus on the parts surrounding the  $i$ -th word of the input sequence.
- Generally, the annotation is a RNN hidden state
- The context vector  $c_i$  is computed as a weighted sum of the annotations  $h_i$ :

$$c_i = \sum_{j=1}^T \alpha_{ij} h_j$$

with  $0 \leq \alpha_{ij} \leq 1$

# Annotation weights

- The weight  $\alpha_{ij}$  of each annotation  $h_j$  is computed by a softmax function

$$\alpha_{ij} = \text{softmax}(e_{ij}) = \frac{\exp(e_{ij})}{\sum_{k=1}^T \exp(e_{ik})}$$

where

$$e_{ij} = \text{score}(s_i, h_j)$$

- The score  $e_{ij}$  measures how well the inputs around position  $j$  and the output at position  $i$  match
- The score is based on the RNN hidden state  $s_i$  (just before emitting  $y_i$ ) and the  $j$ -th annotation  $h_j$  of the input sentence.



# To to compute the score?

- The most popular score are
  - Dot-product:  $e_{ij} = \text{score}(s_i, h_j) = s_i^T h_j = h_j^T s_i$
  - Bilinear function:  $e_{ij} = \text{score}(s_i, h_j) = s_i^T W h_j = h_j^T W s_i$
  - Multi-Layer Perceptron:  $e_{ij} = \text{score}(s_i, h_j) = a(s_i, h_j) = w_2^T \tanh(W_1[s_i, h_j])$
- In the original paper, the authors used the **alignment model**  $a(\cdot)$  as a feedforward neural network which is jointly trained with all the other components of the proposed system.

# Interpretation of $\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^T \exp(e_{ik})}$


- We can understand the approach of taking a weighted sum of all the annotations as computing an **expected annotation**, where the expectation is over possible alignments.
- Let  $\alpha_{ij}$  be a probability that the target word  $y_i$  is aligned to, or translated from, a source word  $x_j$ . Then, the  $i$ -th context vector  $c_i$  is the expected annotation over all the annotations  $h_j$  with probabilities  $\alpha_{ij}$

$$c_i = \sum_{j=1}^T \alpha_{ij} h_j$$

- The probability  $\alpha_{ij}$ , or its associated score  $e_{ij}$ , reflects the importance of the annotation  $h_j$  with respect to the current decoding state  $s_i$  in deciding the next prediction  $y_i$ .
- Intuitively, this implements a **mechanism of attention** in the decoder.

# Illustration of the alignment

- The  $x$ -axis and  $y$ -axis of the plot correspond to the words in the source sentence (English) and the generated translation (French), respectively.
- Each pixel shows the weight  $\alpha_{ij}$  of the annotation of the  $j$ -th source word for the  $i$ -th target word, in grayscale (0: black, 1: white).



# Single-head attention


# Transformer's Encoder: Principle

- Self-attention
- Queries, keys and values
- Scaling the dot product
- Multi-head attention

“Deep Contextualized Word Representations” in NAACL, 2018


Matthew Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, Luke Zettlemoyer

# Example with a sentence



- The word « licked » is most correlated to « cat » (who?) and « owner » (to whom?)

# Comparison with convolution



- The word « licked » is only correlated to words in a given neighborhood (kernel size)

# Self-attention

- Self-attention is a sequence-to-sequence operation:
  - A sequence of vectors goes in, and a sequence of vectors comes out.
  - Let's call the input vectors  $x_1, x_2, \dots, x_t$  and the corresponding output vectors  $y_1, y_2, \dots, y_t$ .
  - The vectors all have dimension  $d$  (the inputs are embedded with an embedding layer).
- To produce output vector  $y_i$ , the self attention operation simply takes a weighted average over all the input vectors

$$y_i = \sum_{j=1}^t w_{ij} x_j$$

where the positive weights  $w_{ij}$  sum to one over all  $j$ .



# Self-attention: basic operation

- The weight

$$w_{ij} = \text{softmax}(e_{ij}) = \text{softmax}(\text{score}(x_i, x_j))$$

is not a parameter, as in a normal neural net, but it is derived from a function over  $x_i$  and  $x_j$ .


- The simplest option for the score function is the dot product:

$$e_{ij} = x_i^T x_j$$

- The dot product gives us a value anywhere between negative and positive infinity, so we apply a softmax to map the values to  $[0,1]$  and to ensure that they sum to 1 over the whole sequence:

$$w_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^t \exp(e_{ik})}$$

# Self-attention



A visual illustration of basic self-attention.  
Note that the softmax operation over the weights is not illustrated.

# Example: How to draw a rose

Query

how to draw a rose

Vidéos :

How To Draw A Rose

YouTube · Art for Kids Hub  
3 févr. 2022

8 temps forts dans cette vidéo

How to Draw a Rose

YouTube · Draw So Cute  
7 févr. 2022

10 temps forts dans cette vidéo

How to Draw + Color a Rose Super EASY Realistic

YouTube · Draw So Cute  
31 janv. 2017

Value

Key

# Queries, keys and values

$$y_i = \sum_{j=1}^t w_{ij} \mathbf{x}_j, \quad w_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^t \exp(e_{ik})}, \quad e_{ij} = \mathbf{x}_i^T \mathbf{x}_j$$

- Every input vector  $\mathbf{x}_i$  is used in three different ways in the self attention operation (each role has a name: query, key or value):

$$y_i = \sum_{j=1}^t \frac{\exp(\mathbf{x}_i^T \mathbf{x}_j)}{\sum_{k=1}^t \exp(\mathbf{x}_i^T \mathbf{x}_k)} \mathbf{x}_j$$

1. **Query**: vector from which the attention is looking for its own output  $y_i$
  2. **Key**: It is compared to every other vector at which the query looks to establish the weights
  3. **Value**: It is used as part of the weighted sum to compute each output vector once the weights have been established.
- In the basic self-attention we've seen so far, each input vector must play all three roles.
  - In the transformer, new vectors for each role are derived, by applying a linear transformation to the original input vector

# Linear transformation for each role

- We can add three  $d \times d$  weight matrices  $W^Q$ ,  $W^K$ ,  $W^V$  to compute three linear transformations of each  $x_i$ , for the three different parts of the self attention:

$$q_i = W^Q x_i, \quad k_i = W^K x_i, \quad v_i = W^V x_i$$

$$e_{ij} = q_i^T k_j$$

$$w_{ij} = \text{softmax}(e_{ij})$$


$$y_i = \sum_{j=1}^t w_{ij} v_j$$

- This gives the self-attention layer some controllable parameters, and allows it to modify the incoming vectors to suit the three roles they must play.

# Attention function: matrix form

- In practice, we compute the attention function on a set of queries simultaneously, packed together into a matrix  $Q$  (after the linear transformation if we use it).
  - Initial vectors are the rows of  $Q$
- The keys and values are also packed together into matrices  $K$  and  $V$ .
- We compute the matrix of outputs as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{Q K^T}{\sqrt{d}}\right) V$$



# Scaling the dot product: why $\sqrt{d}$ ?

- The softmax function can be sensitive to very large input values.
- These kill the gradient, and slow down learning, or cause it to stop altogether.
- Since the average value of the dot product grows with the embedding dimension  $k$ , it helps to scale the dot product back a little to stop the inputs to the softmax function from growing too large:

$$e_{ij} = \frac{q_i^T k_j}{\sqrt{d}}$$

- Why  $\sqrt{d}$ ?
  - Imagine a vector in  $\mathbb{R}^d$  with values all  $c$ :  $(c, c, \dots, c)$ . Its Euclidean length is  $c\sqrt{d}$ .
  - Therefore, we are dividing out the amount by which the increase in dimension increases the length of the average vectors.
- An other theoretical justification of  $\sqrt{d}$  :
  - if all the elements  $q_i$  and  $k_j$  are drawn independently from a  $\mathcal{N}(0, \sigma^2)$  then  $q_i^T k_j$  would have a variance  $d\sigma^4$ .
  - But  $e_{ij}$  has variance of  $\sigma^4$ .
  - This normalization ensures that the numbers given to softmax are not too dispersed.

# Multi-head attention



# Multi-head attention

- Finally, we must account for the fact that a word can mean different things to different neighbours.
  - Consider the following example: « Mary gave roses to Susan »
  - We see that the word gave has different relations to different parts of the sentence.
    - Mary expresses who's doing the giving,
    - roses expresses what's being given,
    - and Susan expresses who the recipient is.
- In a single self-attention operation, all this information just gets summed together.
  - If « Susan gave Mary the roses » instead, the output vector  $y_{gave}$  would be the same, even though the meaning has changed.

# Multi-head attention

- We can give the self attention greater power of discrimination, by combining several self attention mechanisms (which we'll index with  $i$ ), each with different matrices  $W_i^Q$ ,  $W_i^K$ ,  $W_i^V$ . These are called attention heads.
- For input  $x_j$  each attention head produces a different output vector  $y_j^i$ . We concatenate these, and pass them through a linear transformation  $W^O$  to reduce the dimension back to  $d$ .

# Multi-Head Attention

- Just concatenate all the heads and apply an output projection.

$$\text{head}_i = \text{Attention}(W_i^Q x, W_i^K x, W_i^V x)$$
$$\text{MultiHeadedAttention}(x) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O$$


- Previously, we used the following dimensions for **single**-head SA:

$$W^Q \in \mathbb{R}^{d \times d}, \quad W^K \in \mathbb{R}^{d \times d}, \quad W^V \in \mathbb{R}^{d \times d},$$

- In practice, we use a reduced dimension for each head.

$$W_i^Q \in \mathbb{R}^{d \times \frac{d}{h}}, \quad W_i^K \in \mathbb{R}^{d \times \frac{d}{h}}, \quad W_i^V \in \mathbb{R}^{d \times \frac{d}{h}}, \quad W^O \in \mathbb{R}^{d \times d}$$

- The total computational cost is similar to that of single-head attention with full dimensionality.





$h$ : number of heads  
 $d$ : feature dimension  
in output of SA

# Combine with FFN

- Add a **feed-forward network** to add more expressivity.
  - This applies another nonlinearity to the representations (or “post-process” them).

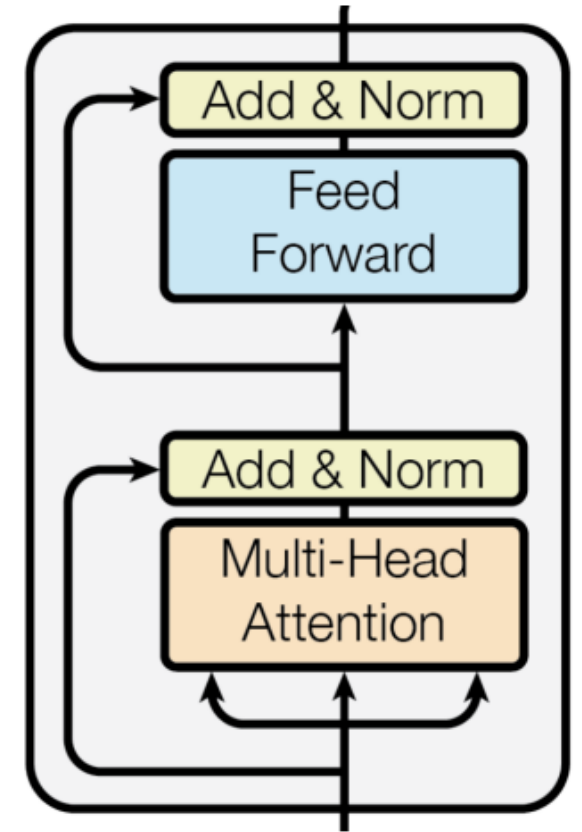
$$\text{FFN}(x) = \sigma(x\mathbf{W}_1 + b_1)\mathbf{W}_2 + b_2$$
$$\mathbf{W}_1 \in \mathbb{R}^{d \times d_{\text{ff}}},$$
$$\mathbf{W}_2 \in \mathbb{R}^{d_{\text{ff}} \times d}$$

- Usually, the dimensionality of the hidden feedforward layer  $d_{\text{ff}}$  is 2-8 times larger than the input dimension  $d$ .



# How Do We Prevent Vanishing Gradients?

- **Residual connections** let the model “skip” layers
  - These connections are particularly useful for training deep networks
- Use **layer normalization** to stabilize the network and allow for proper gradient flow

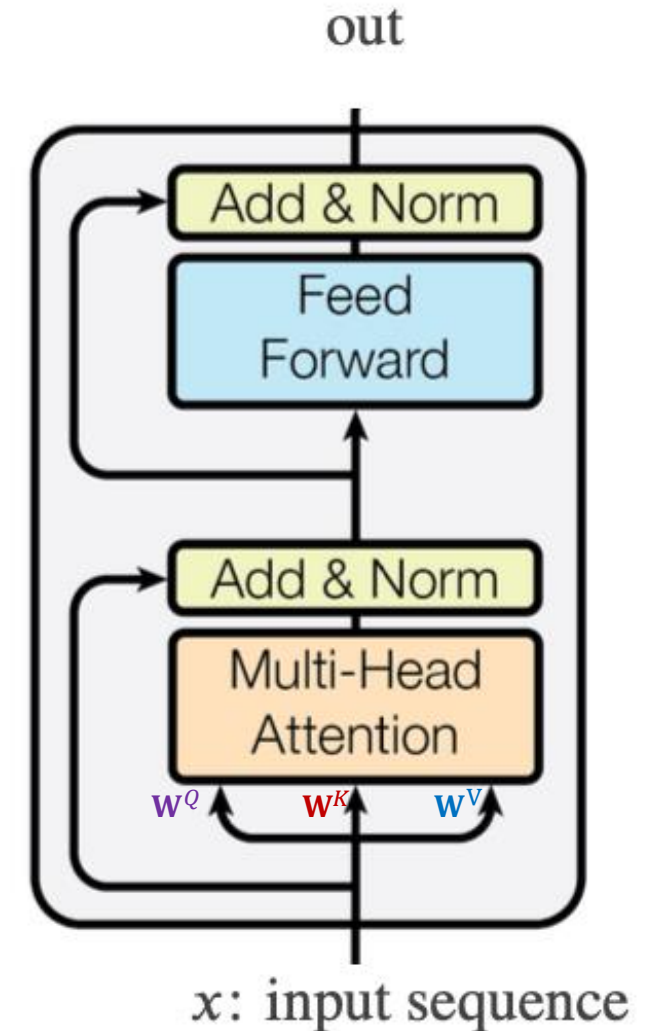


# Putting it Together: Self-Attention Block

- Layer normalization is similar to batch normalization (but it is not strictly speaking a batch normalization)
- Layer normalization prevents the range of values in the layers from changing too much, which allows faster training and better generalization ability
- Given input  $x$ :

$$\begin{aligned} \mathbf{z} &= \text{LayerNorm}(\tilde{\mathbf{x}} + \mathbf{x})\tilde{\mathbf{x}} \\ &= \text{MultiHeadedAttention}(\mathbf{x}; \mathbf{W}^Q, \mathbf{W}^K, \mathbf{W}^V) \end{aligned}$$

$$\begin{aligned} \tilde{\mathbf{z}} &= \text{FFN}(\mathbf{z}) = \sigma(\mathbf{z}W_1 + b_1)W_2 + b_2 \\ \text{out} &= \text{LayerNorm}(\tilde{\mathbf{z}} + \mathbf{z}) \end{aligned}$$



# LayerNorm( $x + \text{Sublayer}(x)$ )

- For a batch  $\{x_n\}_{n=1,\dots,N}$  of  $N$  vectors  $x_n \in \mathbb{R}^K$ , also written as  $\{x_{n,k}\} \in \mathbb{R}^{N \times K}$ , the expectation and variance accros spatial dimensions are « estimated » by

$$\mu_n = \frac{1}{K} \sum_{k=1}^K x_{n,k} \in \mathbb{R}, \quad \sigma_n^2 = \frac{1}{K} \sum_{k=1}^K (x_{n,k} - \mu_n)^2 \in \mathbb{R}$$

- Layer Normalization (LayerNorm in Pytorch)

$$\hat{x}_{n,k} = \frac{x_{n,k} - \mu_n}{\sqrt{\sigma_n^2 + \epsilon}} \in \mathbb{R} \quad \Rightarrow \quad \hat{x}_n = \begin{pmatrix} \hat{x}_{n,1} \\ \vdots \\ \hat{x}_{n,K} \end{pmatrix} \in \mathbb{R}^K$$

$$LN_{\gamma, \beta}(x_n) = \gamma \hat{x}_n + \beta$$

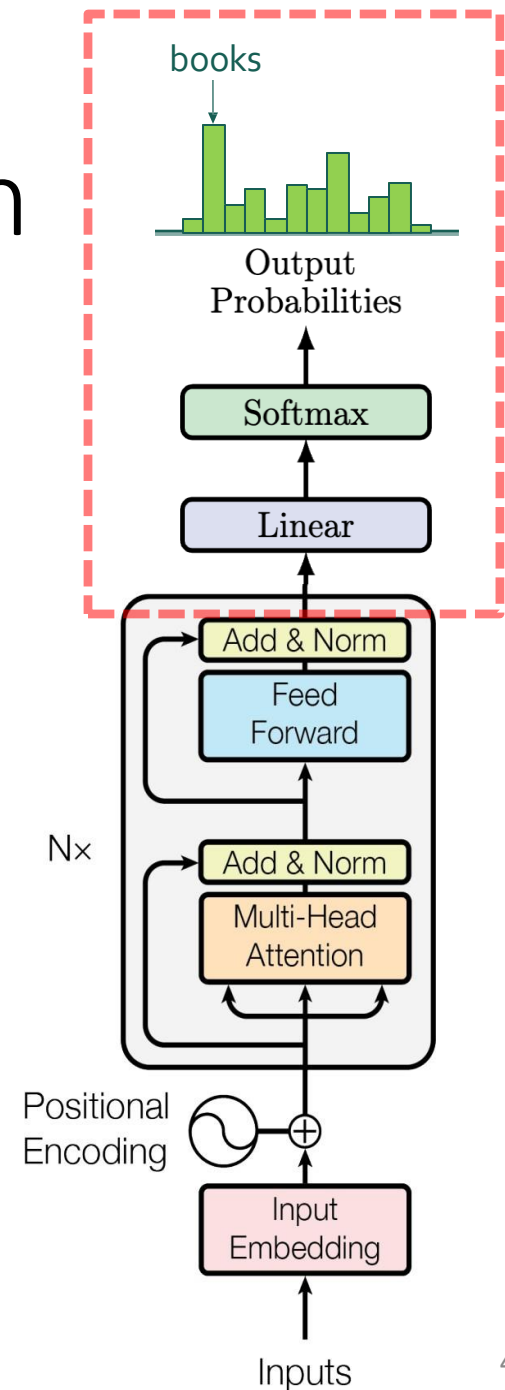
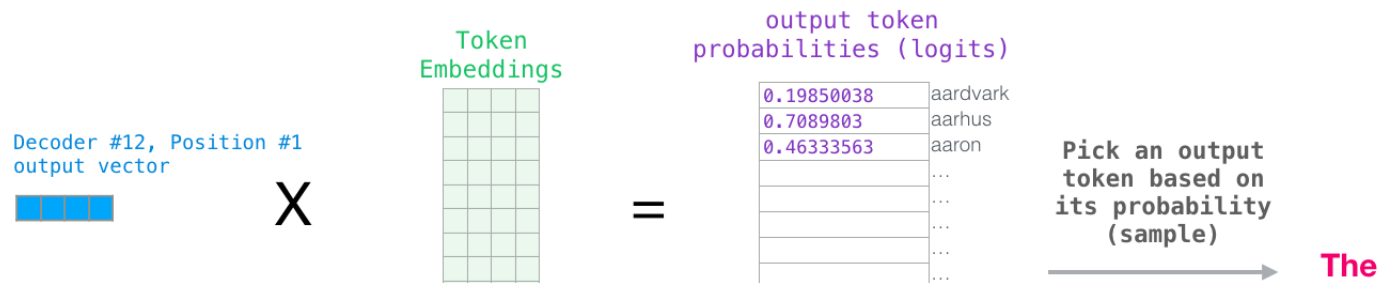
- $\gamma$  and  $\beta$  are learnable affine transform parameters

# Decoder-Only Transformer



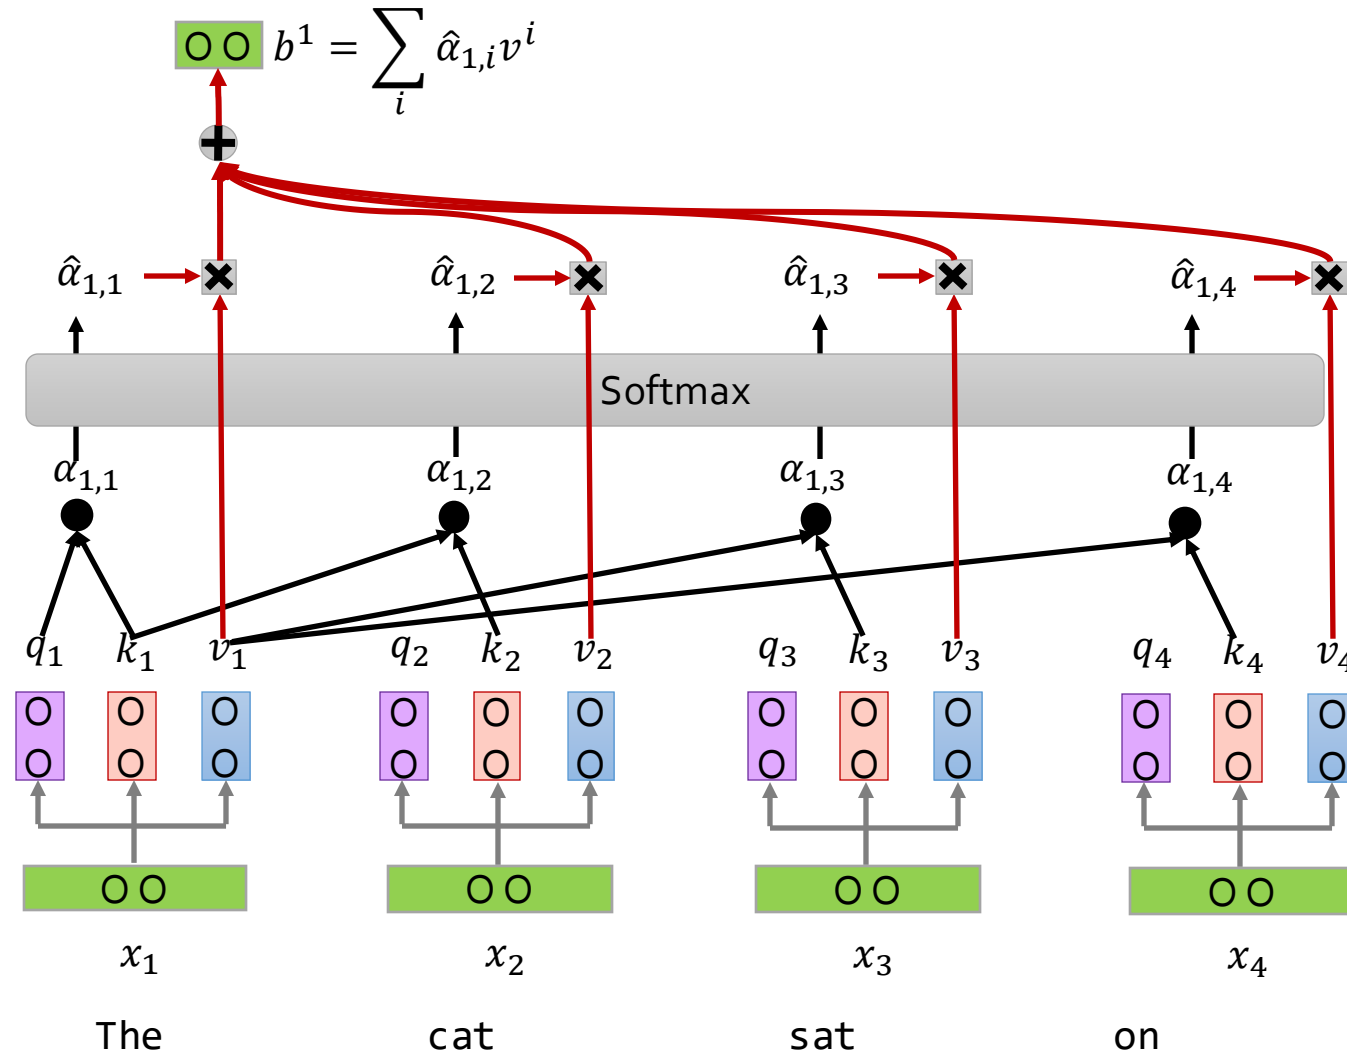
# From Representations to Prediction

- Use sequentially  $N$  Multi-Head attention modules.
- To perform prediction, add a classification head on top of the final layer of the transformer:  $x \in \mathbb{R}^{n \times d}$ .
  - $n$  is the length of the input sequence.
- To obtain logits, we can apply a linear transformation with token embedding matrix  $W^S \in \mathbb{R}^{d \times V}$
- To obtain probabilities, run this through softmax:
 
$$\text{softmax}(xW^S) \in \mathbb{R}^{n \times V}$$

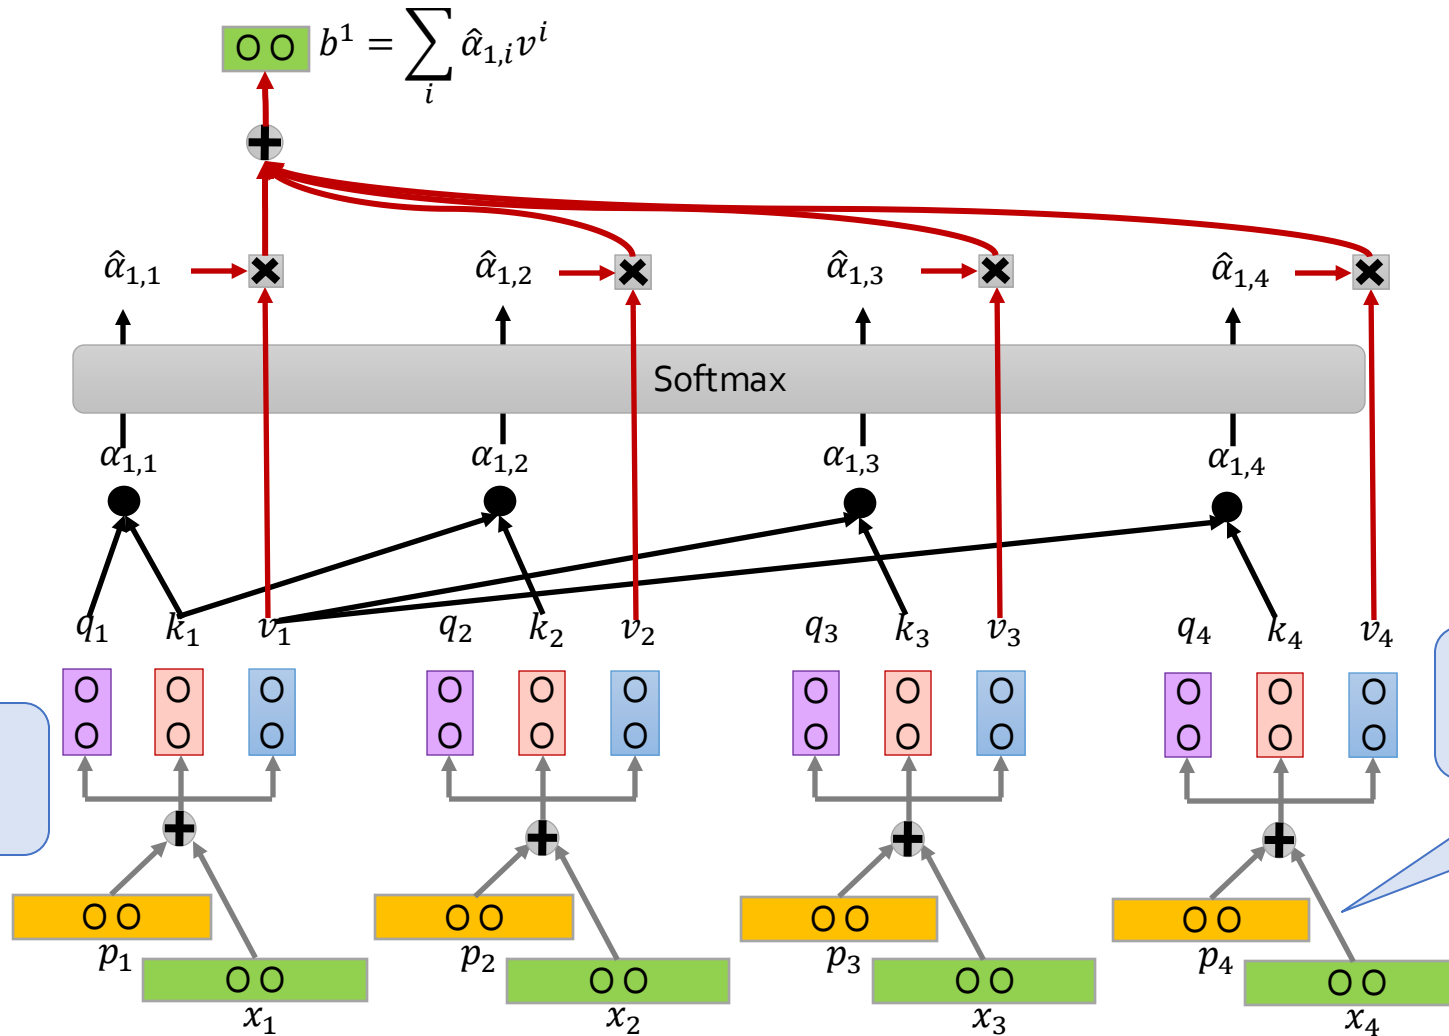


# Input encoding

# Positional Embeddings



# Positional Embeddings



Allows model to learn about positions

# Criteria for Positional Encodings

- The first idea that might come to mind is to assign a number to each time-step within the  $[0, 1]$  range in which 0 means the first word and 1 is the last time-step.
  - One of the problems it will introduce is that you can't figure out how many words are present within a specific range. In other words, time-step doesn't have consistent meaning across different sentences.
- Another idea is to assign a number to each time-step linearly: the first word is given "1", the second word is given "2", and so on.
  - The problem with this approach is that not only the values could get quite large, but also our model can face sentences longer than the ones in training.
  - In addition, our model may not see any sample with one specific length which would hurt generalization of our model.
- Ideally, the following criteria should be satisfied:
  - It should output a unique encoding for each time-step (word's position in a sentence)
  - Distance between any two time-steps should be consistent across sentences with different lengths.
  - The model should generalize to longer sentences without any efforts. Its values should be bounded.
  - It must be deterministic.

# Proposed Method for Transformer

- The encoding proposed by the authors satisfies all of those criteria.
  - First of all, it isn't a single number. Instead, it's a  $d$ -dimensional (same dimension as word embedding) vector  $p_t$  that contains information about a specific position  $t$  in a sentence.
  - Secondly, this encoding is not integrated into the model itself. Instead, this vector is used to equip each word with information about its position in a sentence.
  - According to the authors, for any fixed offset  $s$ ,  $p_{t+s}$  can be represented as a linear function of  $p_t$
- Let  $t$  the desired position in an input sentence,  $p_t = (p_t(0), \dots, p_t(d-1)) \in \mathbb{R}^d$  be its corresponding encoding. Then,


$$p_t(i) = \begin{cases} \sin(w_k t) & \text{if } i = 2k \\ \cos(w_k t) & \text{if } i = 2k + 1 \end{cases} \quad \text{with} \quad w_k = \frac{1}{10000^{2k/d}}$$

- Example:  $p_t(0) = \sin\left(\frac{t}{10000^{0/d}}\right)$ ,  $p_t(1) = \cos\left(\frac{t}{10000^{0/d}}\right)$ ,  $p_t(2) = \sin\left(\frac{t}{10000^{2/d}}\right)$ ,  $p_t(3) = \cos\left(\frac{t}{10000^{2/d}}\right)$ , ...

# Visualizing the Positional Encodings

- The 128-dimensional positional encoding for a sentence with the maximum length of 50.
- Each row represents the embedding vector  $p_t$

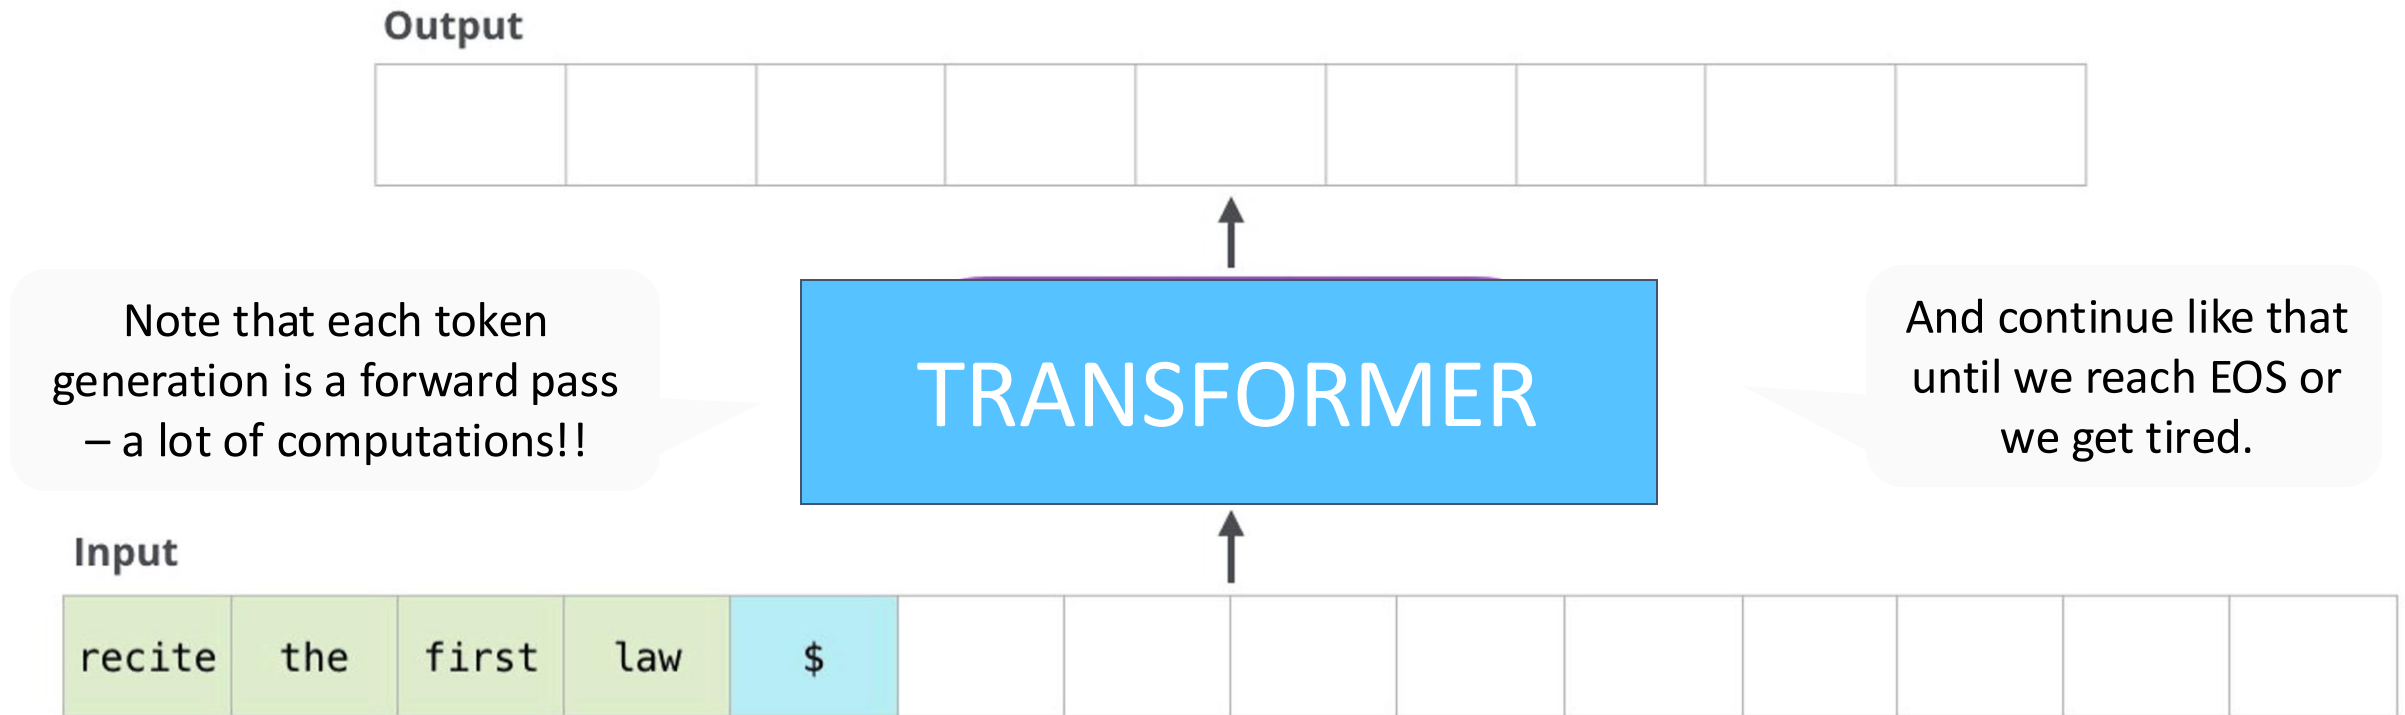
Any pair of rows are different!



How does training work?



# Generating text via Transformer




# Training a Transformer Language Model

- **Goal:** Train a Transformer for language modeling (i.e., predicting the next word).
- **Approach:** Train it so that each position is predictor of the next (right) token.
  - We just shift the input to right by one, and use as labels

$Y =$  cat sat on the mat </s>

EOS special token

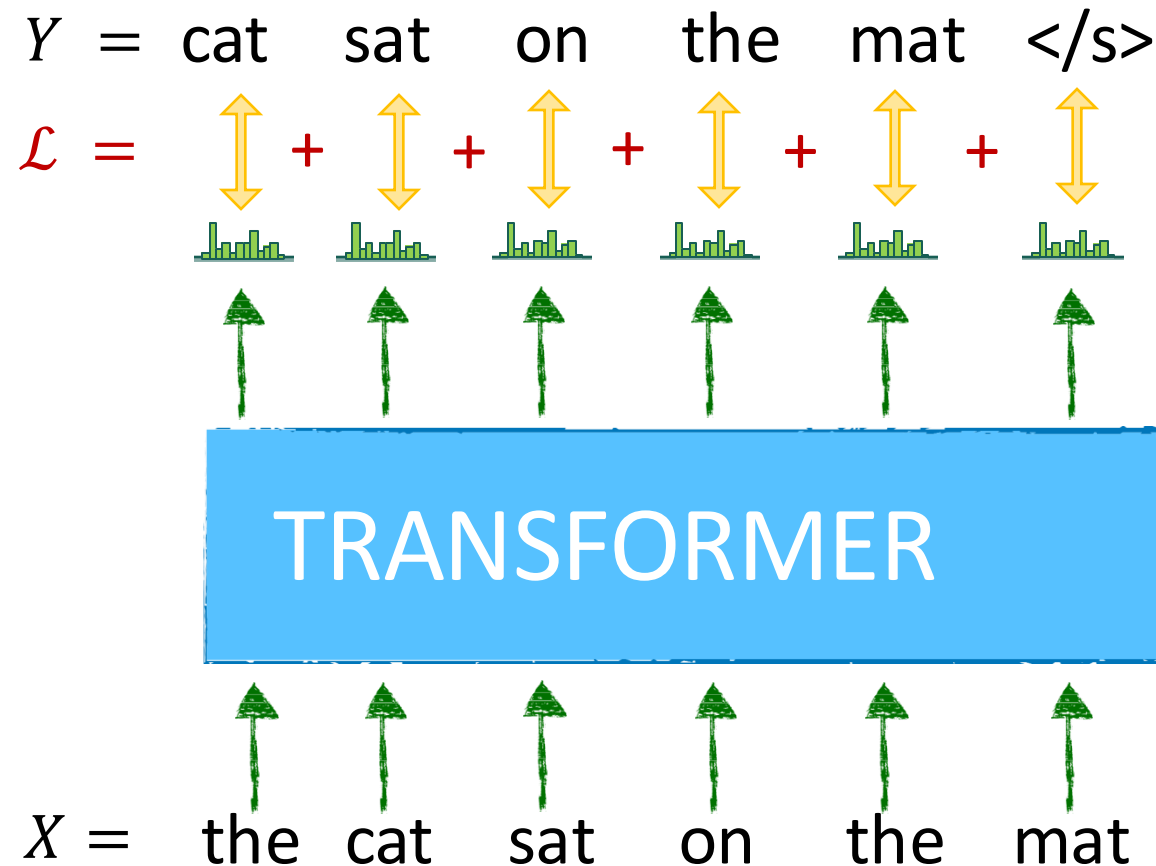


```
X = text[:, :-1]  
Y = text[:, 1:]
```

$X =$  the cat sat on the mat

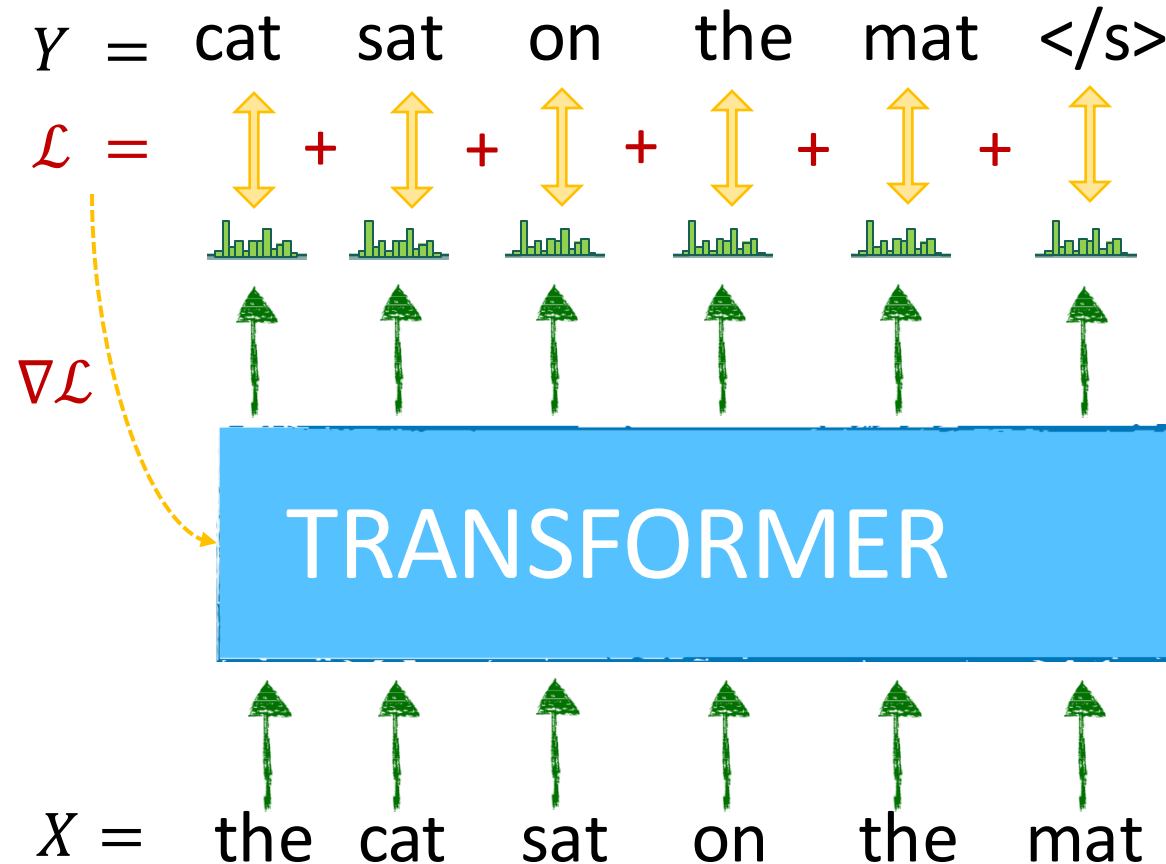
# Training a Transformer Language Model

- Sum the position-wise loss values to obtain a **global loss**.



# Training a Transformer Language Model



- The model would solve the task by **copying** the next token to output (data leakage) since we process the input sequence as a whole.
- Does **not** learn anything useful




# Attention mask

Attention raw scores

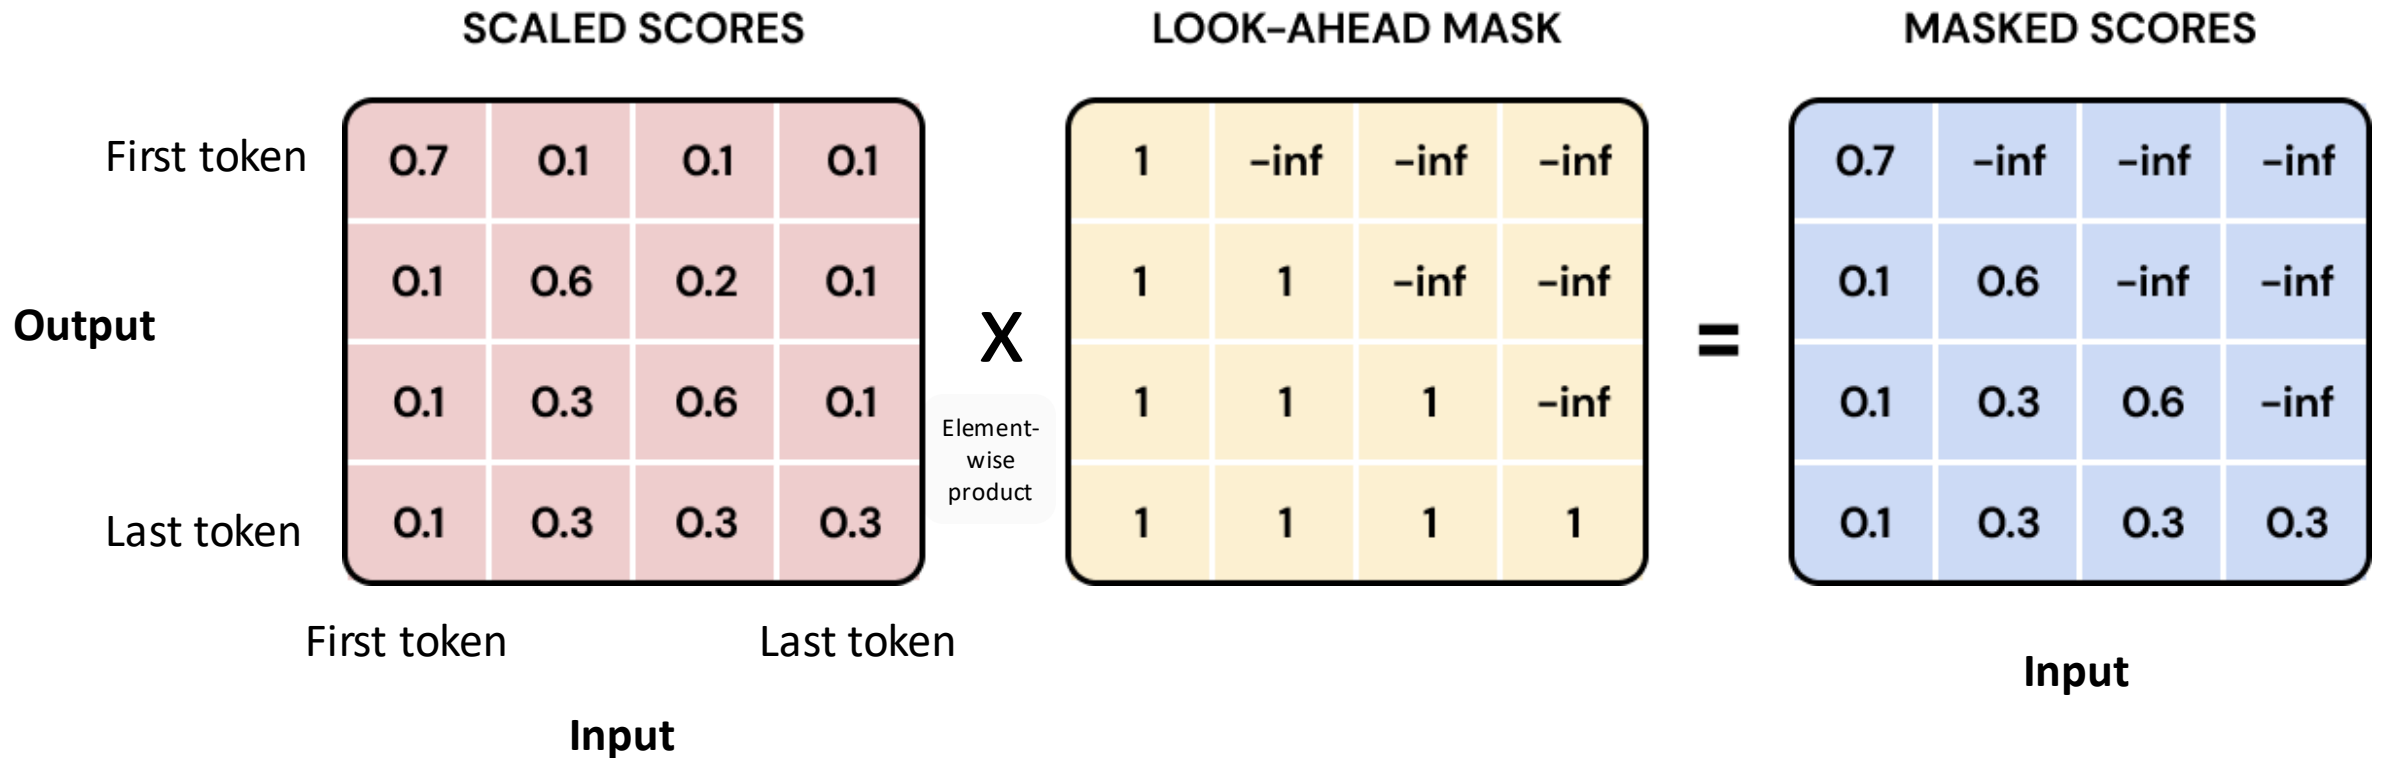
0	-0.08	1.24	0.69	-0.98	1.43	-0.6	0.7	0.16	0.93	1.28	-1.61	-1.1
1	-0.09	-0.0	-0.7	0.06	0.25	0.23	0.26	0.18	0.78	-0.21	-1.01	1.01
2	0.86	1.19	1.59	0.86	-0.13	-0.15	-2.13	-0.98	-0.87	-1.72	1.87	-0.72
3	0.12	-0.03	-0.02	0.88	-0.46	-0.7	0.54	-0.42	-1.89	-0.38	0.04	-0.84
4	0.51	0.17	0.13	-1.64	0.24	-0.02	1.68	-0.36	0.64	0.36	0.27	0.66
5	0.24	-1.44	0.43	0.74	0.96	-1.21	-0.31	1.54	1.66	1.14	0.58	-1.44
6	0.26	-0.1	0.93	0.72	-0.38	1.65	0.47	-0.96	-0.17	-0.9	-1.57	0.22
7	-0.55	0.81	0.71	1.7	-0.8	-1.14	-0.32	1.78	-0.7	-0.04	1.54	0.81
8	0.74	-0.76	-0.44	-0.08	-1.38	-0.13	1.25	-1.37	1.84	0.3	0.57	0.74
9	-0.97	-0.91	0.15	0.35	-0.81	0.11	1.14	-1.52	1.06	1.87	0.5	-0.3
10	1.56	0.9	0.39	1.46	1.44	-1.05	0.9	-0.73	0.36	-0.67	-0.62	-0.43
11	0.32	0.74	0.44	-0.1	1.19	0.83	0.29	2.06	0.51	-0.26	1.51	0.11
	1	2	3	4	5	6	7	8	9	10	11	12



# Attention mask




Attention mask



# Attention masking: Why Before Softmax?

- We applied attention masking before softmax. Why not after?



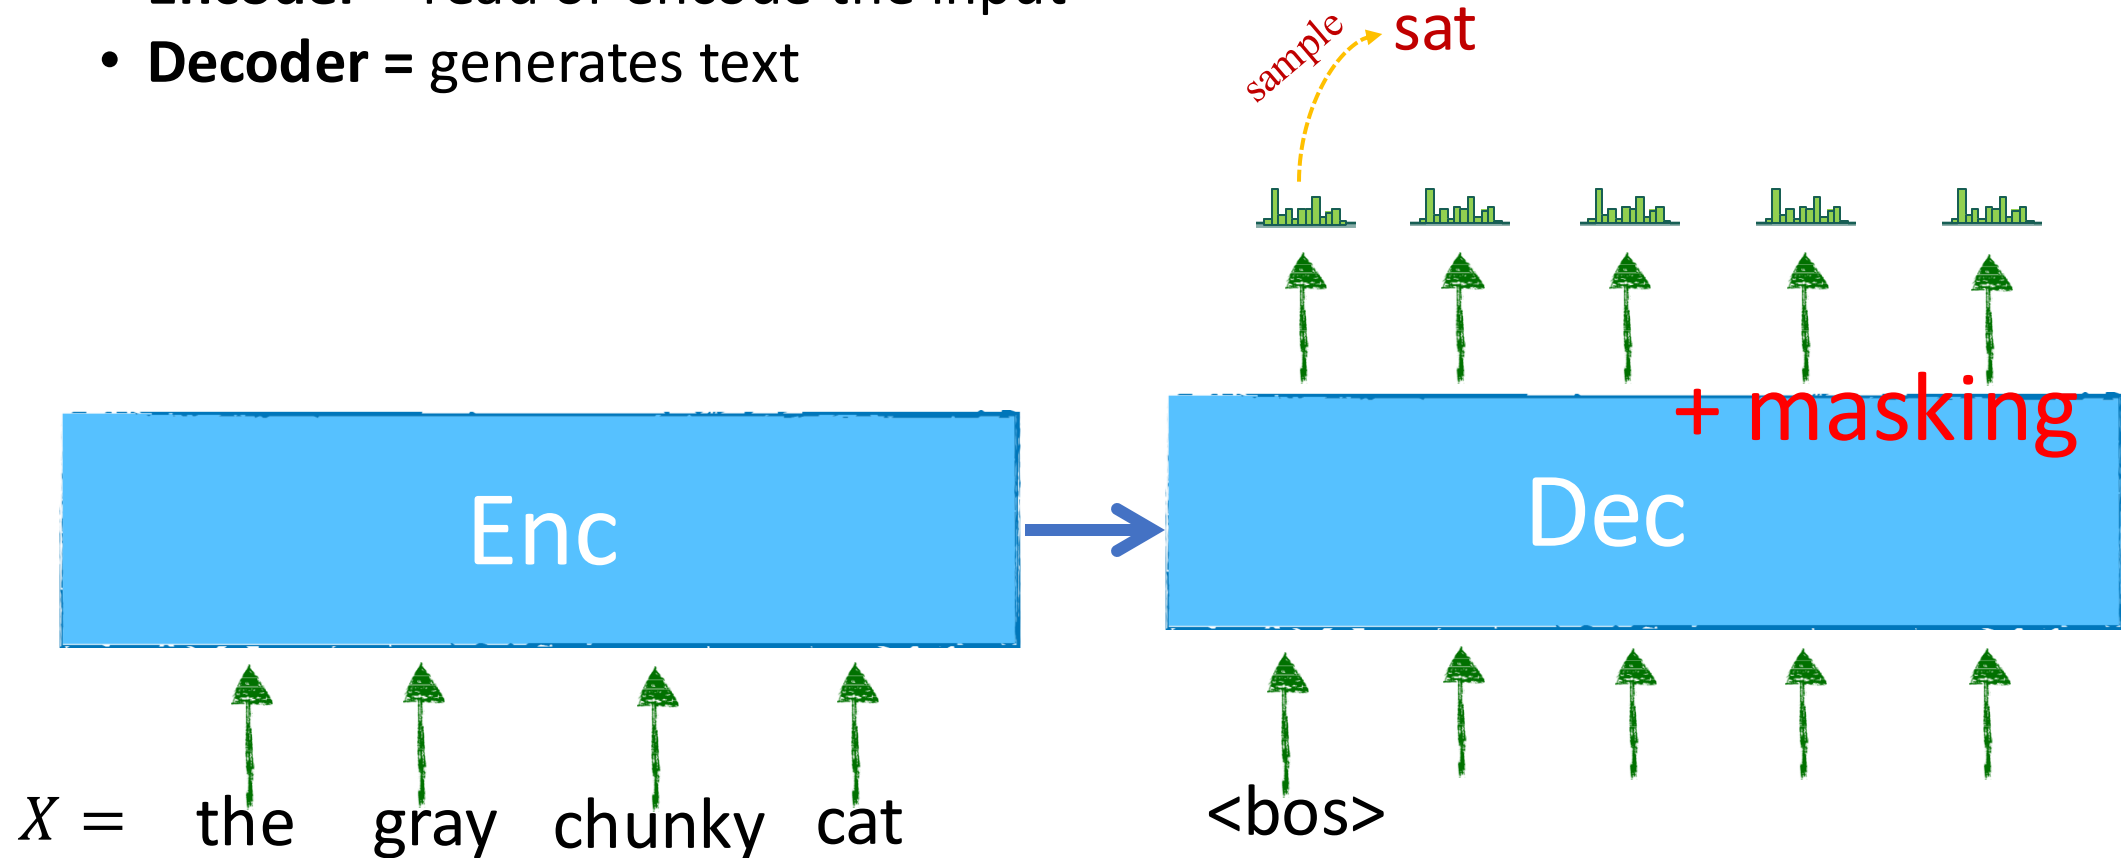
- Softmax normalizes the scores so that it's a probability.
- Masking after softmax, would lead to an unnormalized probability distribution.

# Encoder-Decoder Transformer



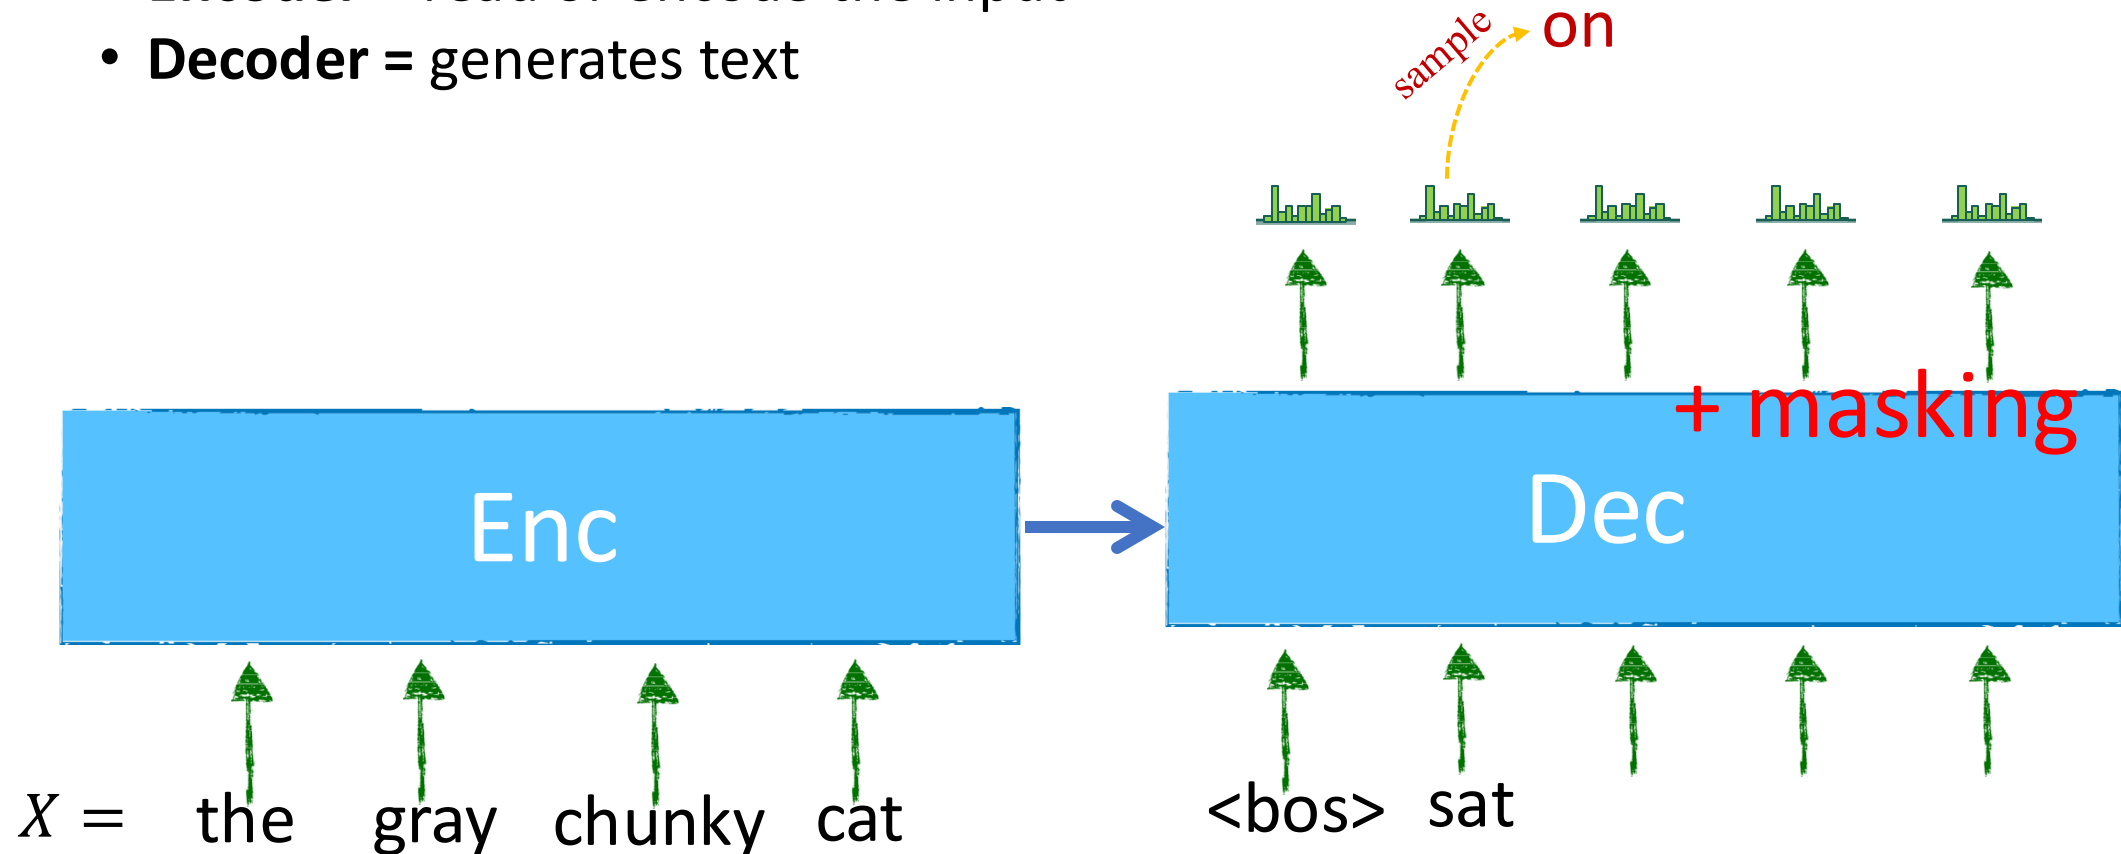
# Enc-Dec work at inference time

- Transformer is two blocks
  - **Encoder** = read or encode the input
  - **Decoder** = generates text



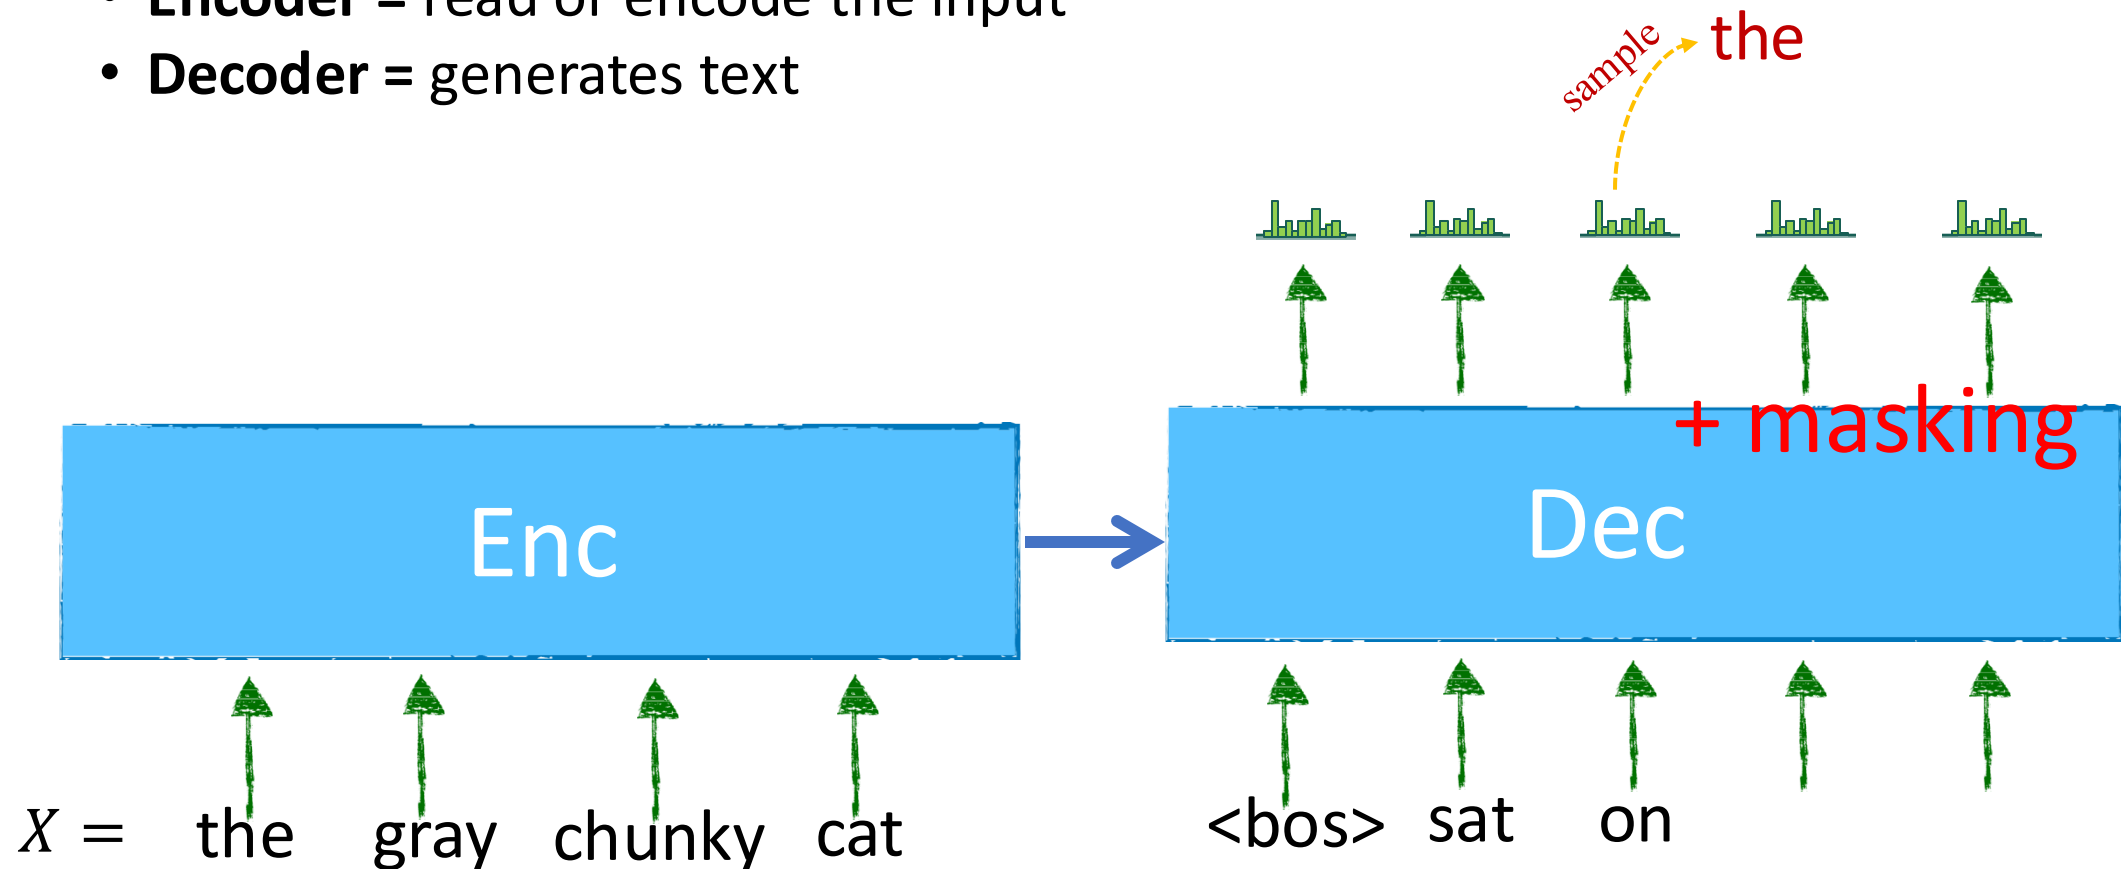
# Enc-Dec work at inference time

- Transformer is two blocks
  - **Encoder** = read or encode the input
  - **Decoder** = generates text



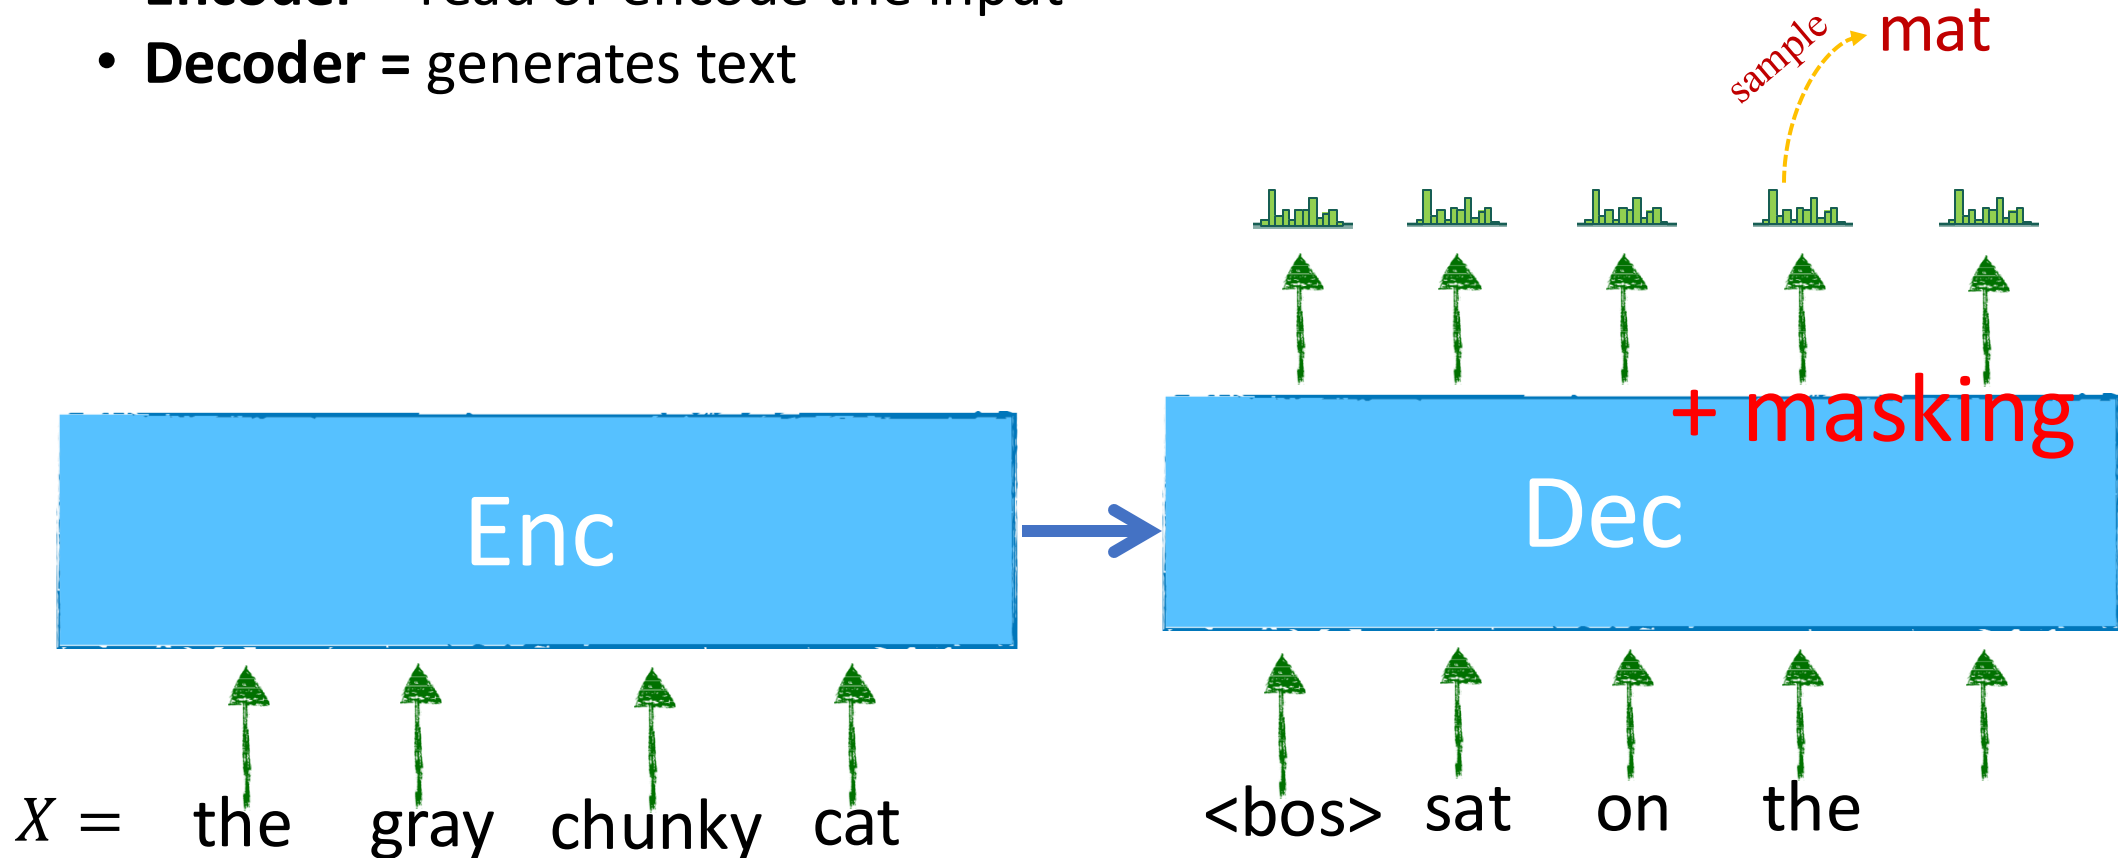
# Enc-Dec work at inference time

- Transformer is two blocks
  - **Encoder** = read or encode the input
  - **Decoder** = generates text




# Enc-Dec work at inference time

- Transformer is two blocks
  - Encoder** = read or encode the input
  - Decoder** = generates text

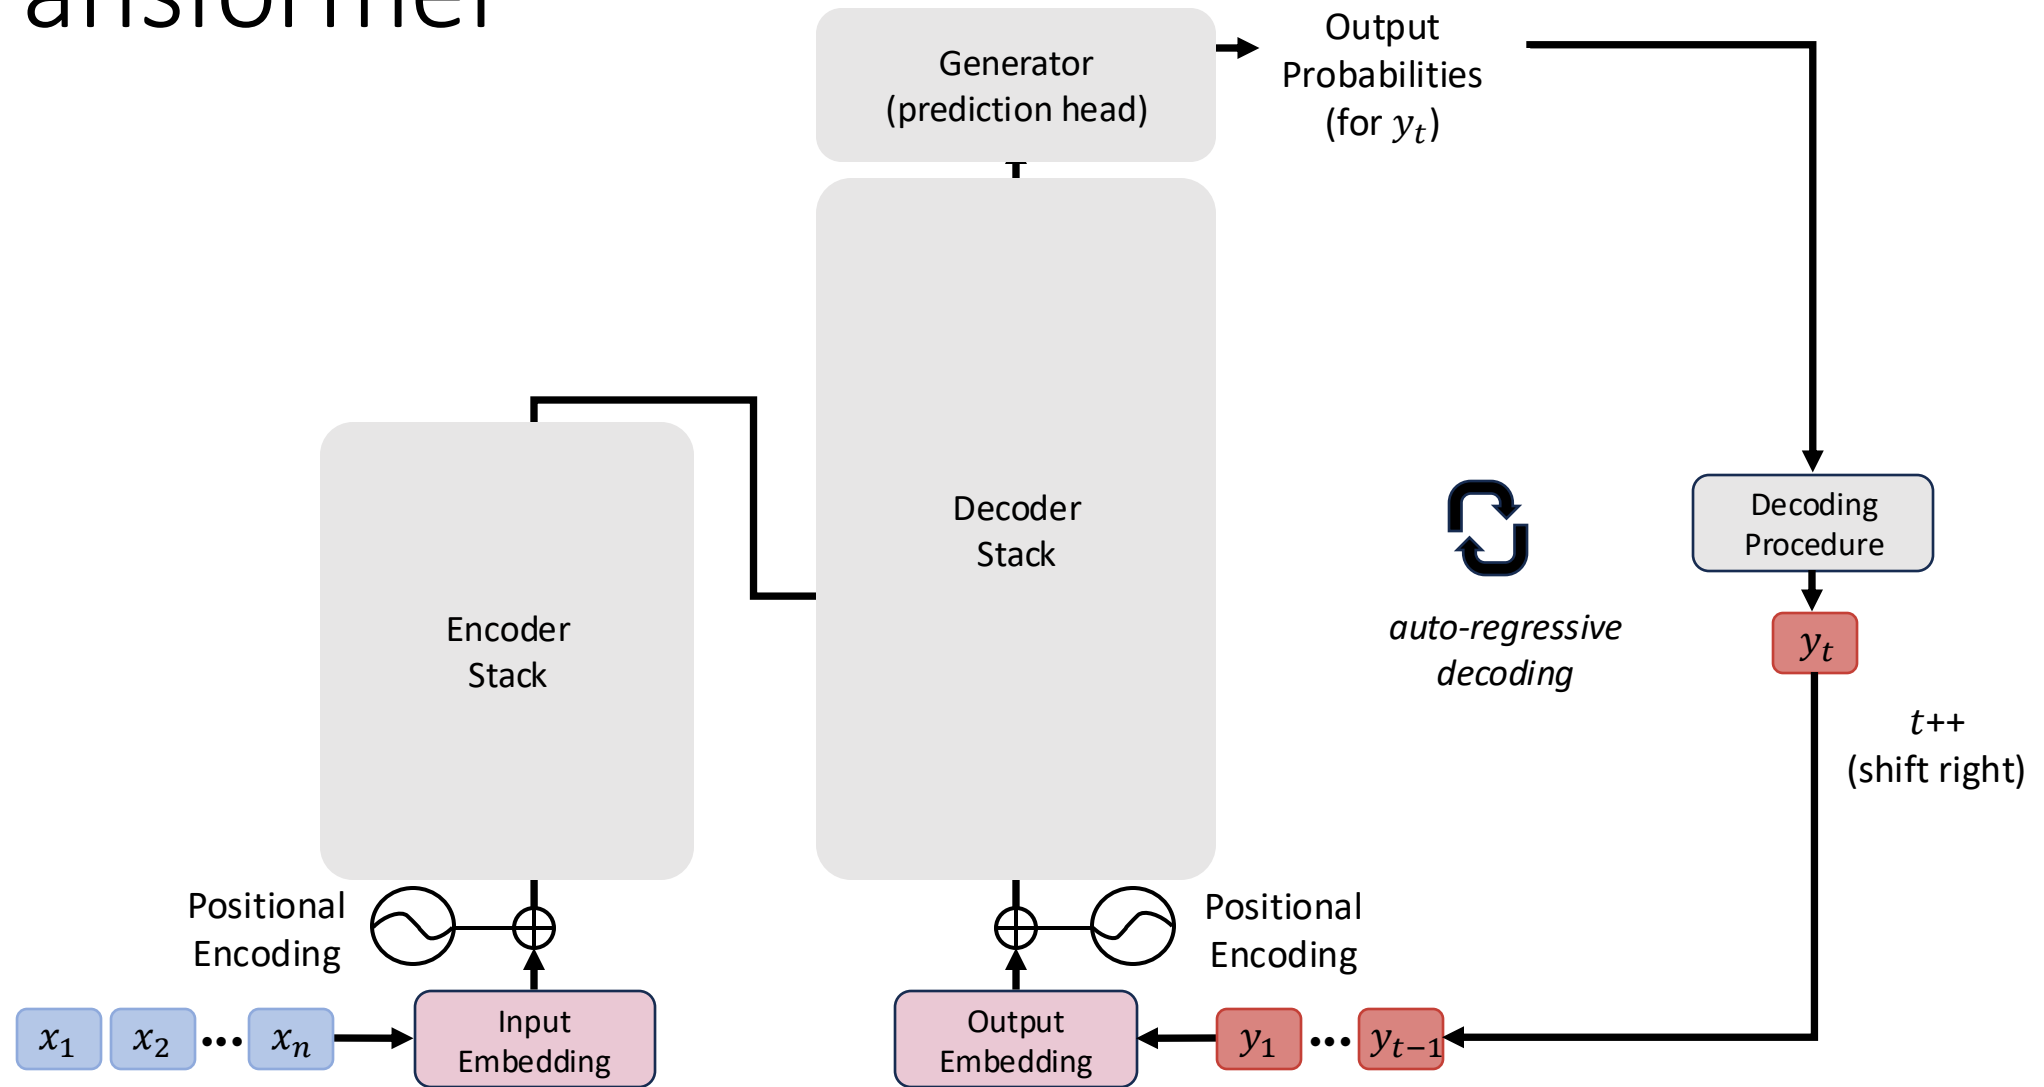


# Enc-Dec work at inference time

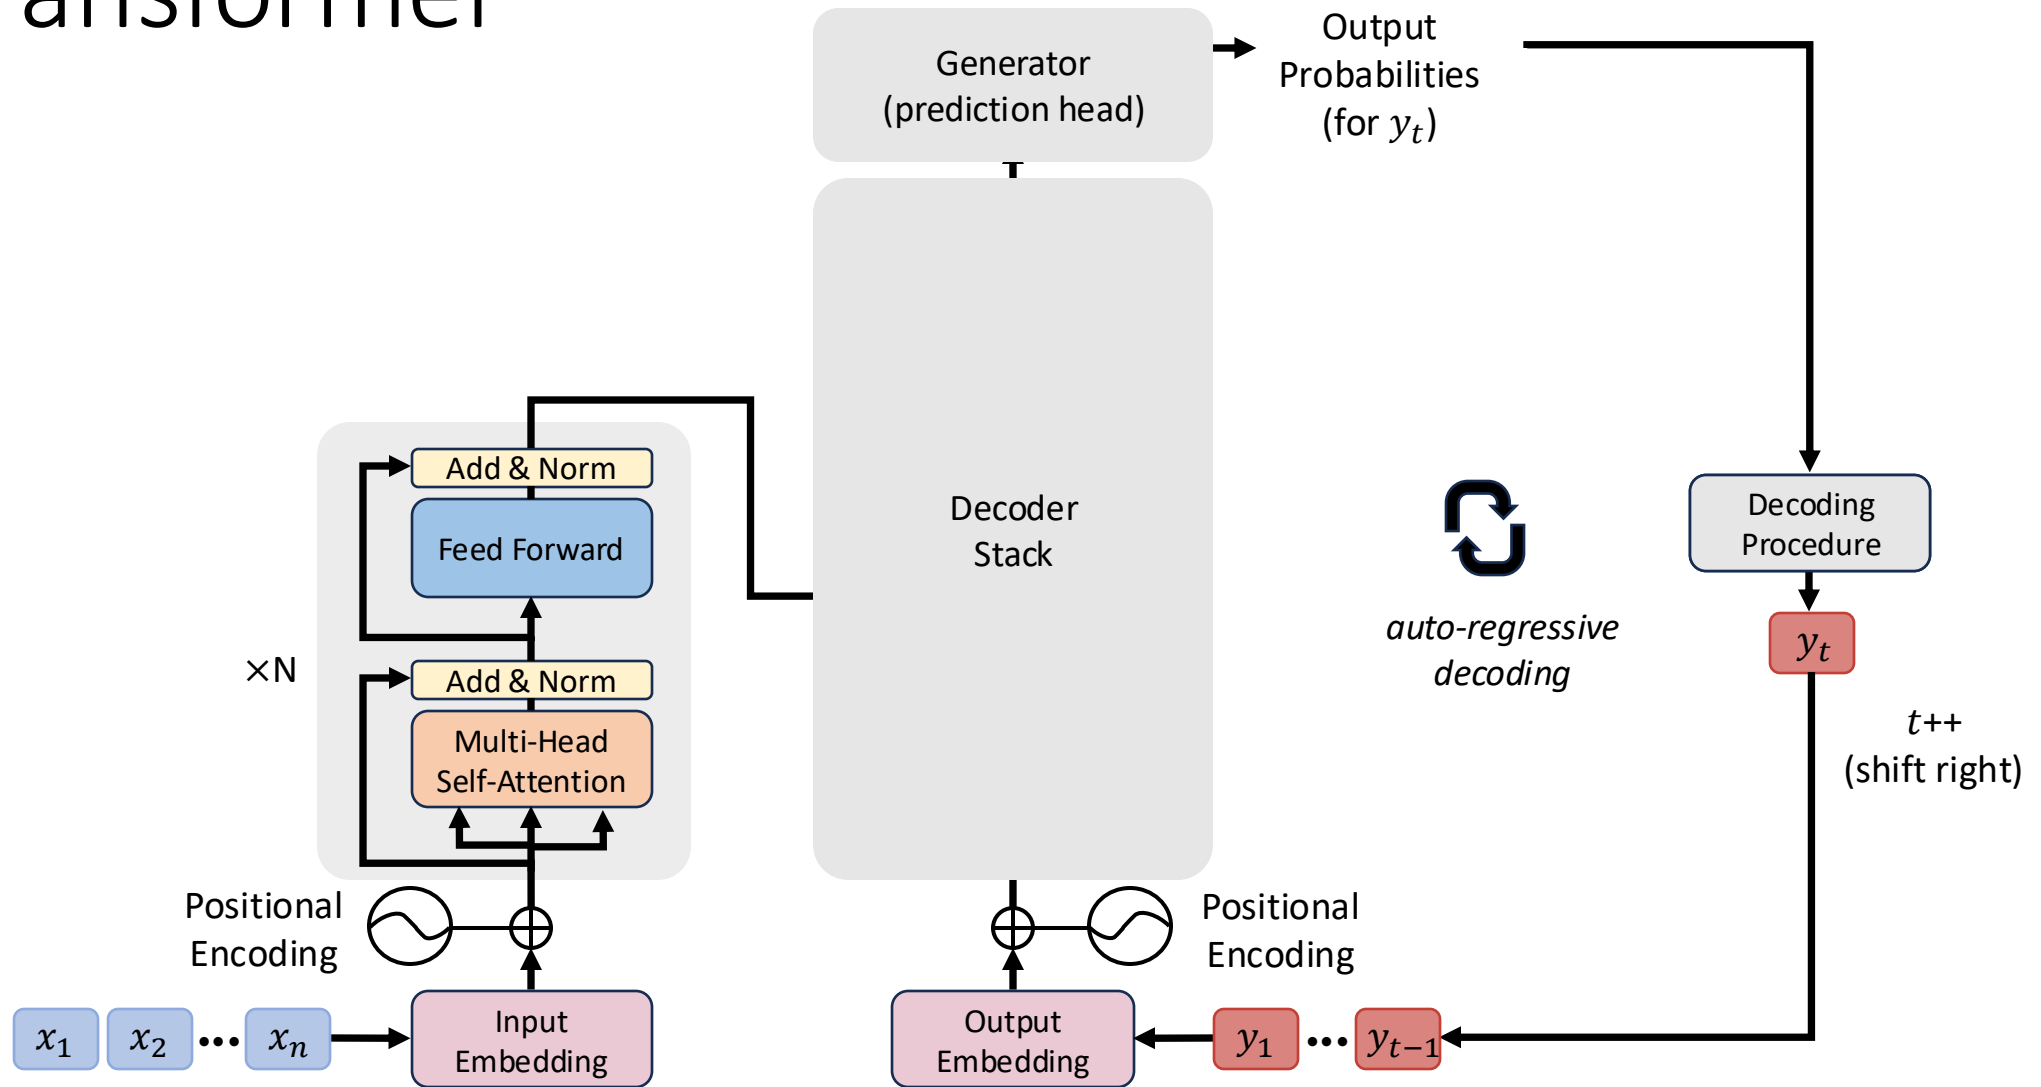
- Transformer is two blocks
  - Encoder** = read or encode the input
  - Decoder** = generates text



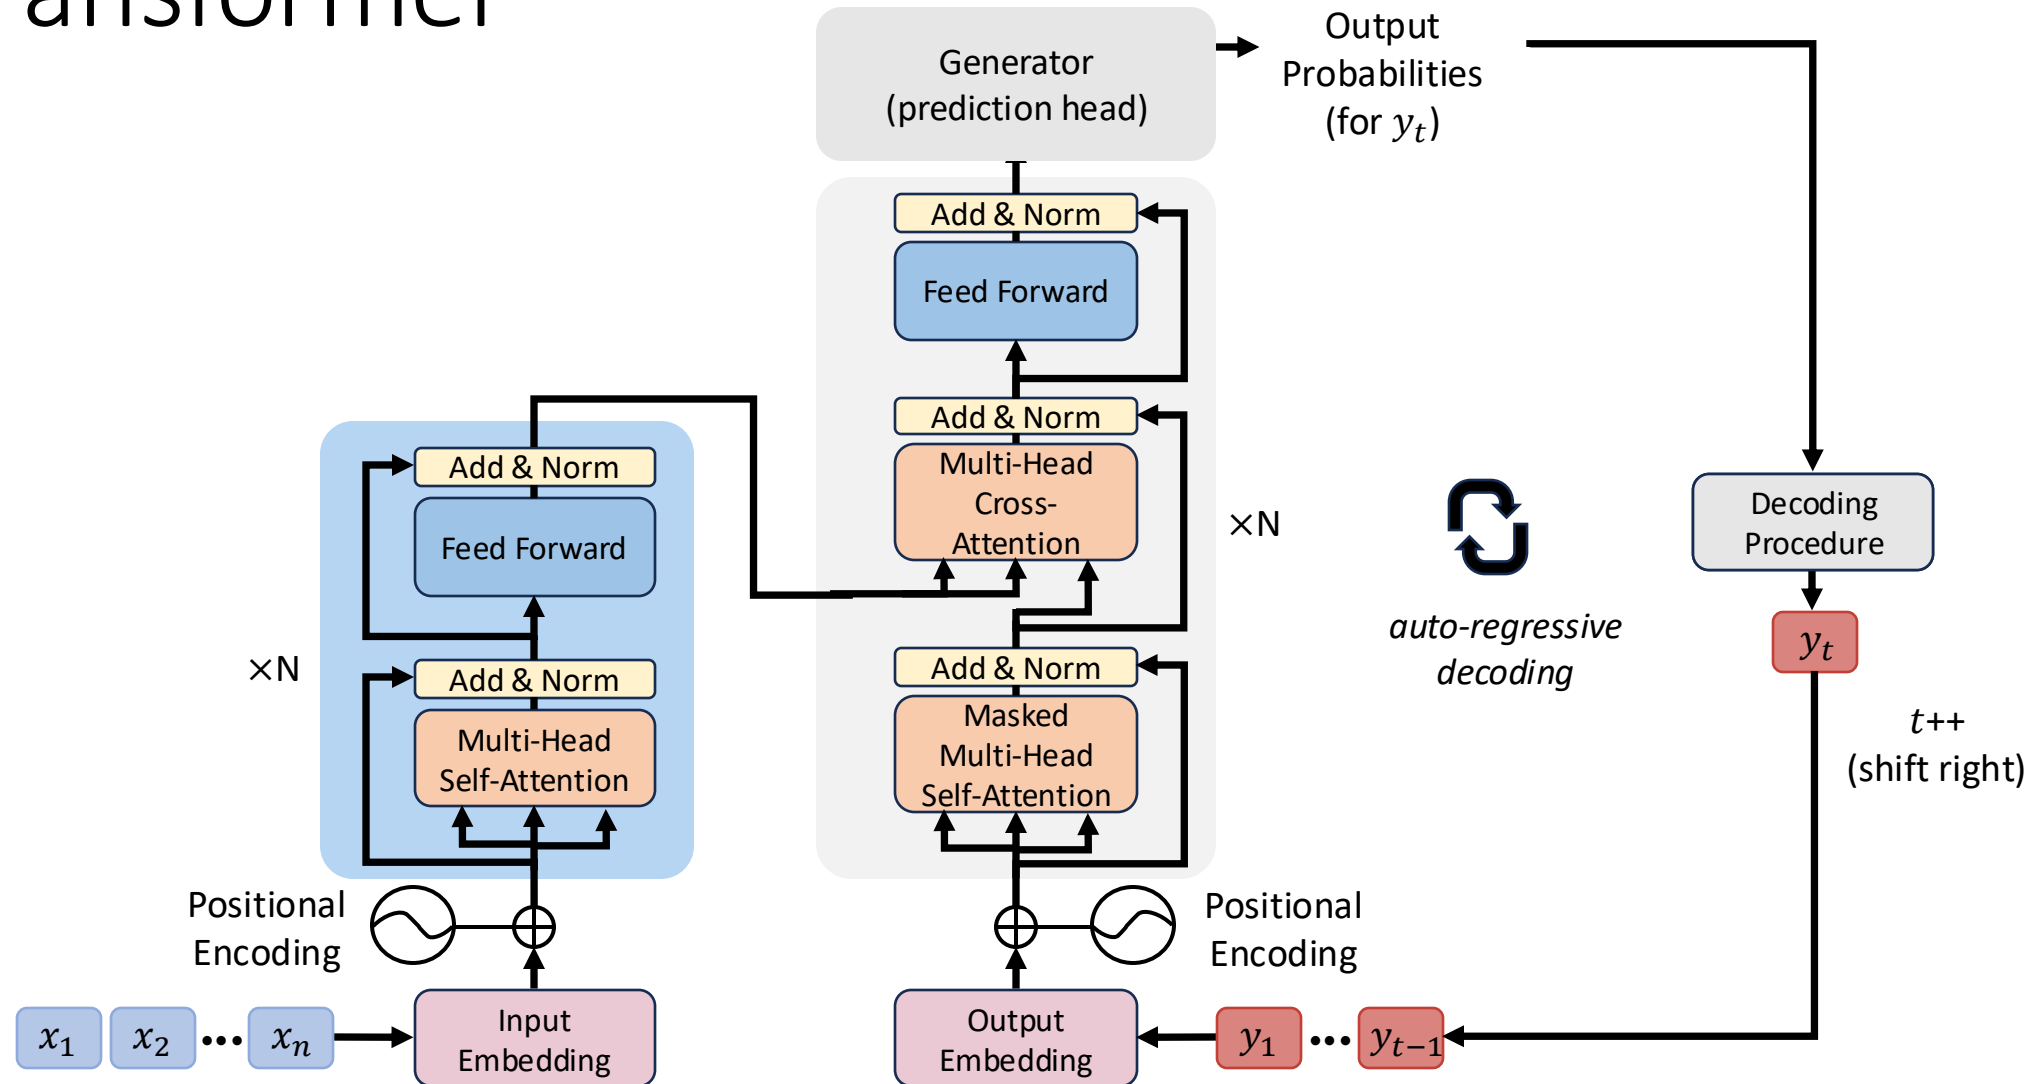
# Transformer



# Transformer

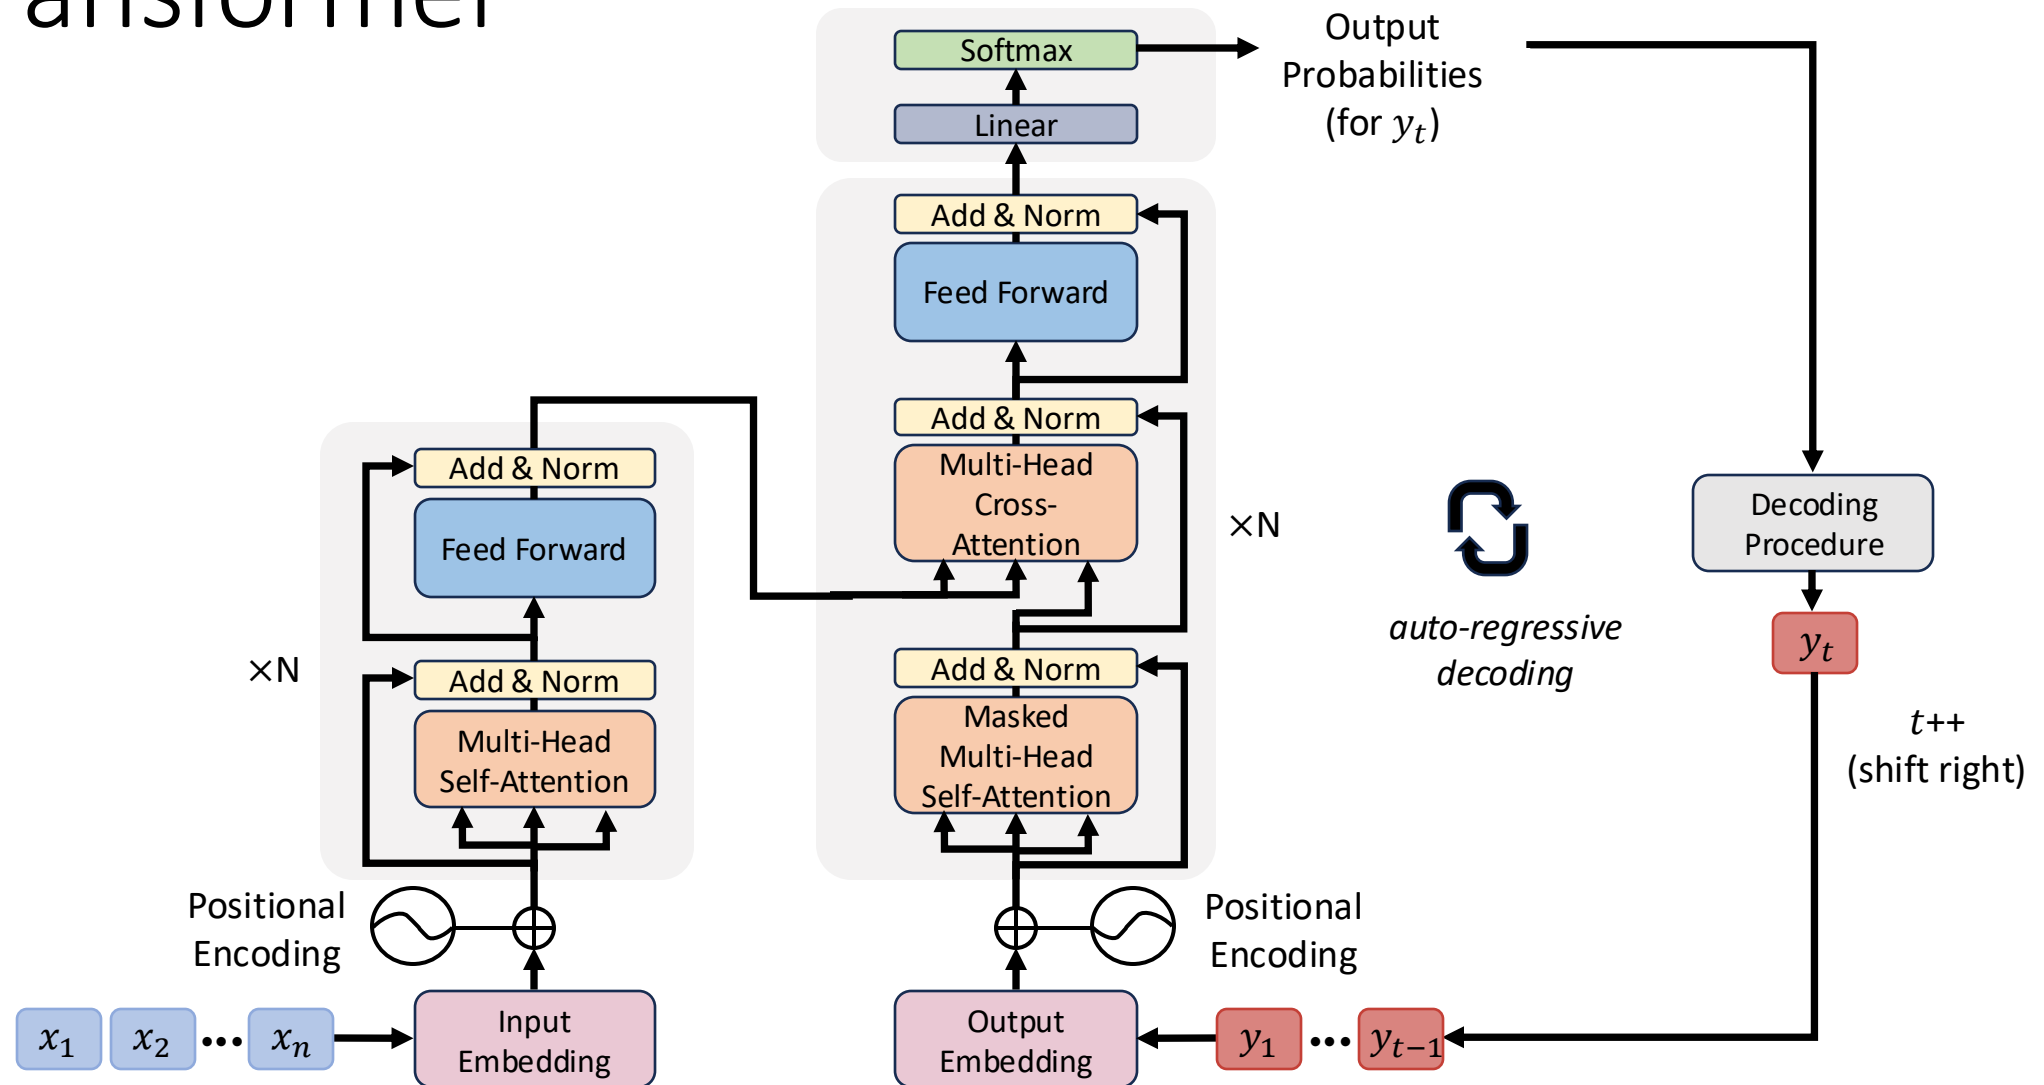


# Transformer





# Transformer

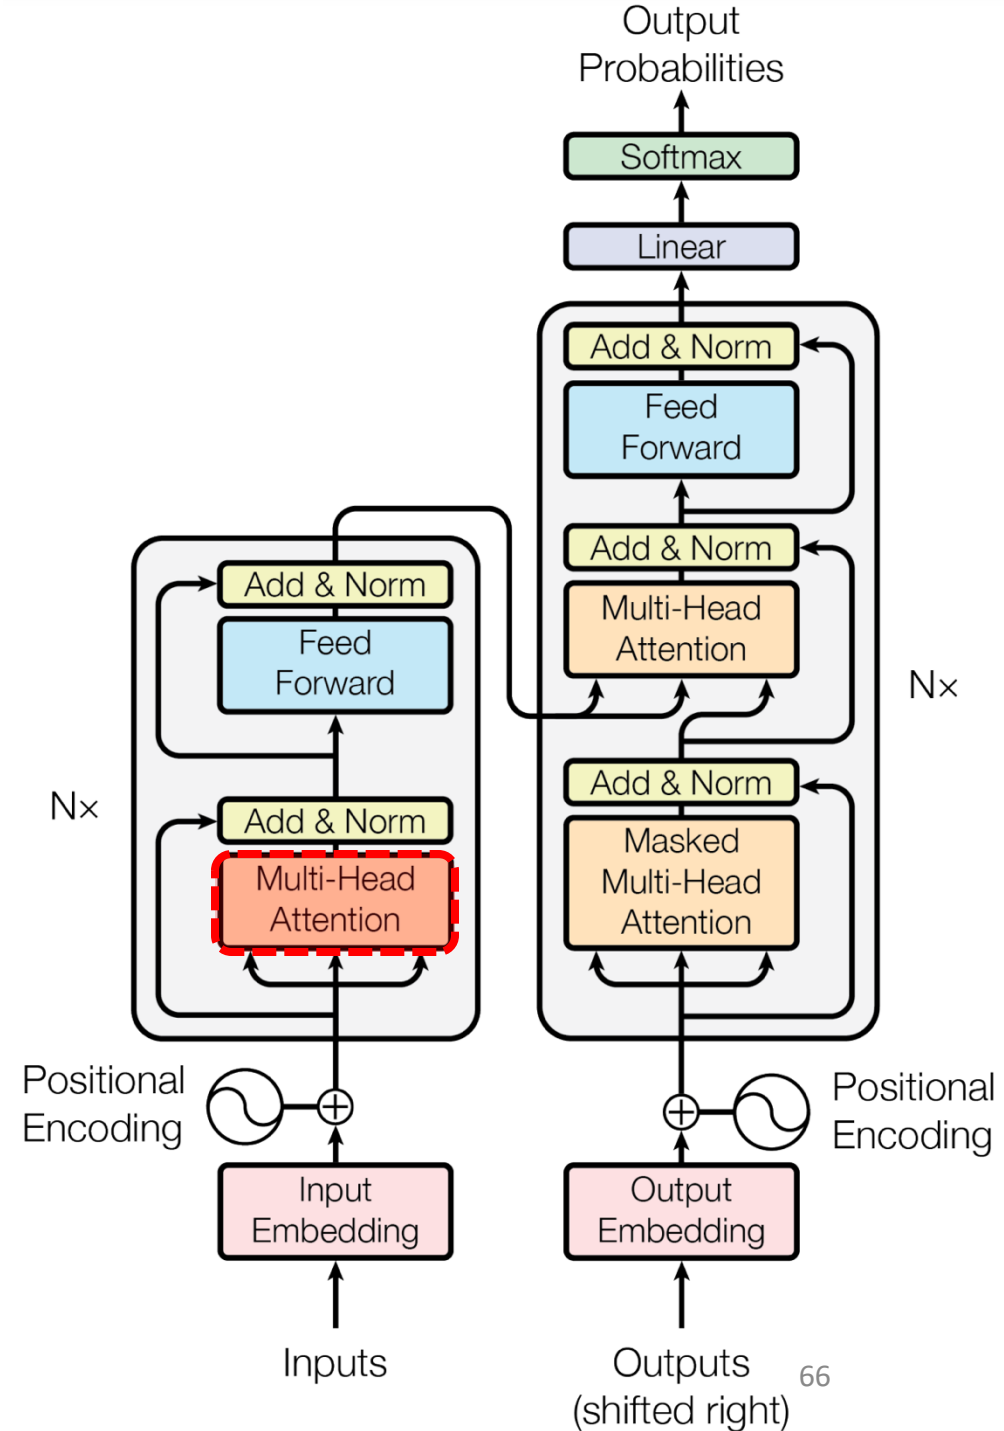



# Transformer

- Computation of **encoder** attends to both sides.

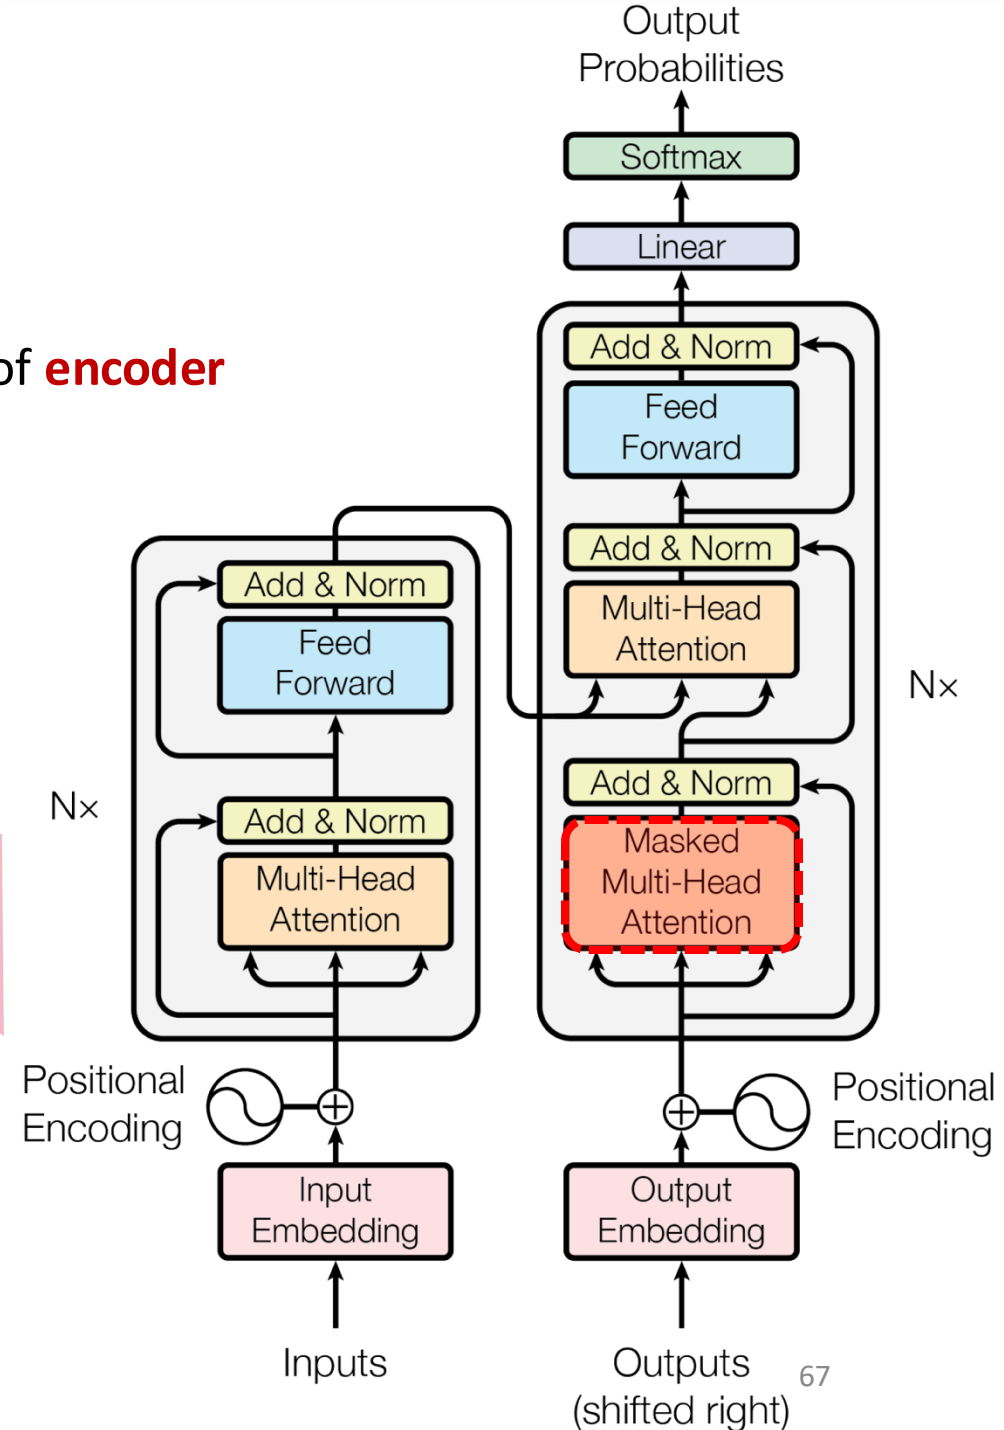
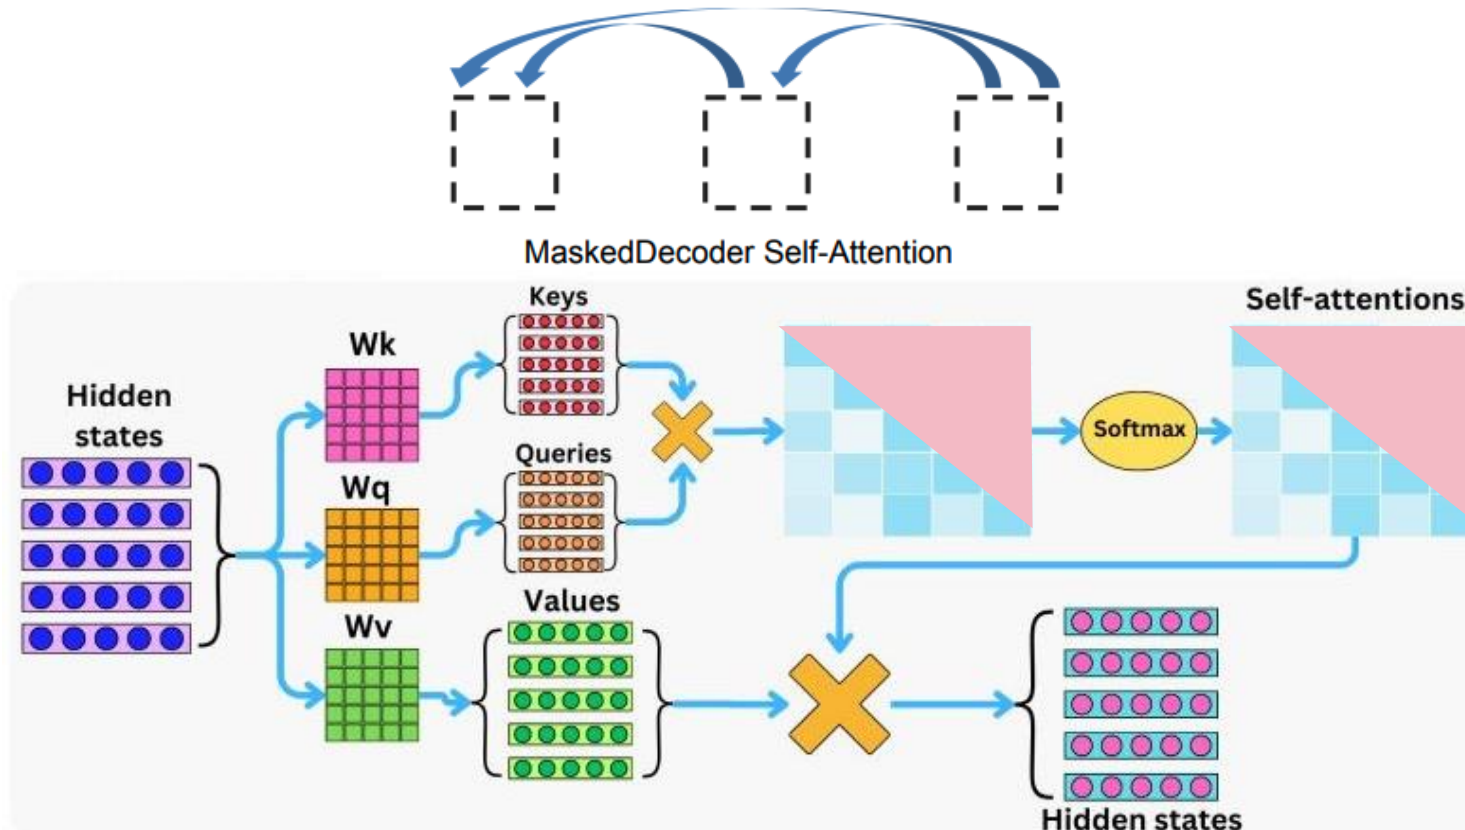


Encoder Self-Attention




# Transformer

- At any step of **decoder**, it attends to previous computation of **encoder** as well as **decoder's** own generations.

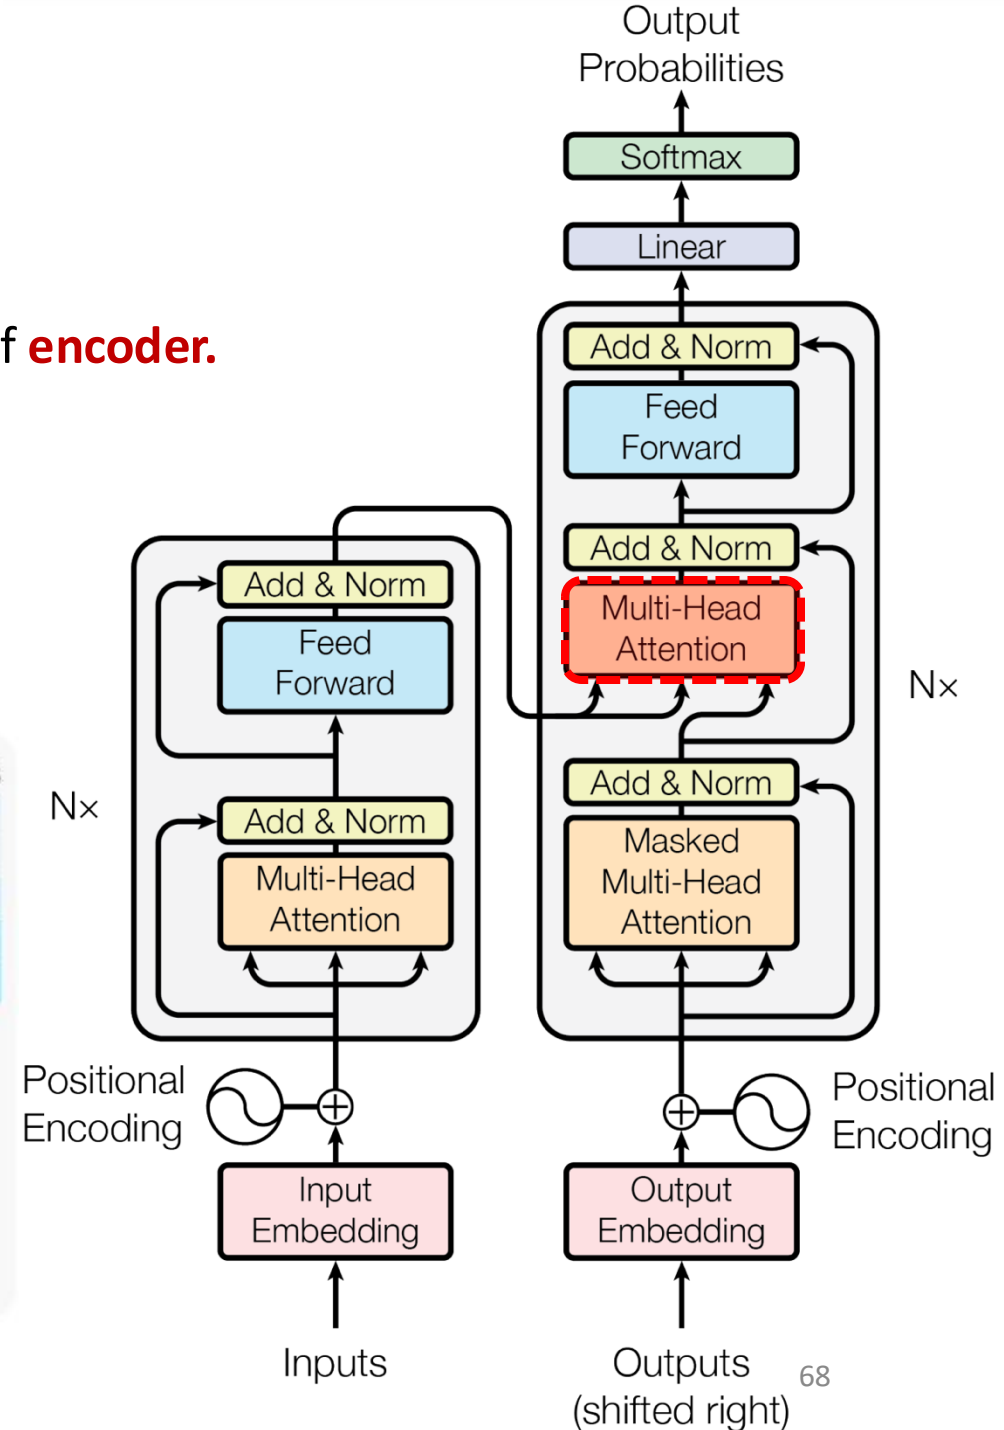



# Transformer

- At any step of **decoder**, it attends to previous computation of **encoder**.



Encoder-Decoder Attention



# Conclusion

# Conclusion

- Attention is nowadays a crucial mechanism for deep neural networks
- Transformers are exploiting self-attention
- Transformers are powerful and generic
- There are many ways to aggregate transformers.
- Each method has advantages and disadvantages.